

Verifying Equivalence of Spark Programs

Shelly Grossman¹, Sara Cohen², Shachar Itzhaky³, Noam Rinetzky¹, and Mooly Sagiv¹

- 1 School of Computer Science, Tel Aviv University, Tel Aviv, Israel
`{shellygr,maon,msagiv}@tau.ac.il`
- 2 School of Engineering and Computer Science, The Hebrew University of Jerusalem, Jerusalem, Israel
`sara@cs.huji.ac.il`
- 3 Computer Science and Artificial Intelligence Laboratory, Massachusetts Institute of Technology, USA
`shachari@mit.edu`

Abstract

In this paper, we present a novel approach for verifying the equivalence of Spark programs. Spark is a popular framework for writing large scale data processing applications. Such frameworks, intended for data-intensive operations, share many similarities with database systems, but do not enjoy a similar support of optimization tools used by traditional databases. Our goal is to enable such optimizations by first providing the necessary theoretical setting for verifying the equivalence of Spark programs. This is challenging because such programs combine relational algebraic operations with *User Defined Functions (UDFs)*.

We model Spark as a programming language which imitates Relational Algebra queries in the bag semantics and allows for user defined functions expressible in Presburger Arithmetics. We present a sound verification technique for verifying equivalence of Spark programs which is complete for programs without aggregate operations as well as for programs with aggregate operations which fall under criteria we provide.

1 Introduction

The rise of Cloud computing and Big Data in the last decade allowed the advent of new programming models, with the intention of simplifying the development process for large-scale needs, letting the programmer focus on business-logic and separating it from the technical details of data management over computer clusters, distribution, communication and parallelization. The first model was MapReduce [10], It allowed programmers to define their logic by composing several iterations of map and reduce operations, where the programmer provided the required map and reduce functions in each step, and the framework was responsible for facilitating the dataflow to the provided functions. MapReduce gave a powerful, yet a clean and abstract programming model. Later on, other frameworks were developed, allowing programmers to write procedural code while still retaining the ability to run it on a large computer cluster, without having the programmer to handle distribution and error recovery. One such framework is Apache Spark [16], in which programmers keep writing code in their programming language of choice, (e.g., Scala [3], Java [1], Python [2], or R [12]) but utilize a special object provided by Spark, called *resilient distributed dataset (RDD)*, providing access to the distributed data itself and to perform transformations on it, using the cloud resources for actual computing. The architecture of Spark comprises of a single master node, referred to as the *driver*, and *worker* nodes in a clustered computer environment. All the nodes have access to the program code, but the driver orchestrates its execution using the underlying Spark framework, abstracting away communications, error recovery, distribution, data partitioning, and parallelization. The access to the data is via



© Shelly Grossman, Sara Cohen, Shachar Itzhaky, Noam Rinetzky, and Mooly Sagiv;
licensed under Creative Commons License CC-BY

Leibniz International Proceedings in Informatics
LIPIC Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl Publishing, Germany

the *RDD* API. The *RDD* can be thought of as a simple database table, but which provides support for *User Defined Functions* (*UDF*), greatly increasing its expressive power. Spark programs handle a family of common tasks involving large datasets, such as log parsing, database queries (via *SparkSQL* [5]), training algorithms and different numeric computations in various fields. Due to this, many Spark programs share several properties: they are mostly short and relatively simple to read and understand. We believe that thanks to this nature of Spark programs, the problem of verifying a program's properties, or even program equivalence as we focus on in this paper, may become feasible, even decidable in a usable class of Spark programs.

Main Results. In this paper we define a simple programming language called SparkLite, in which operations on a single RDD are abstracted as composite simply typed λ -calculus expressions. The operations correspond to operations in relational algebra, with additional aggregate operations. We describe the problem of *program equivalence* (*PE*), and provide a classification of SparkLite programs according to the decidability of the *PE* problem for programs in the same class. When *PE* is decidable, we show an algorithm for solving it.

Overview. Section 2 provides necessary preliminaries for the rest of the paper. In section 3 we give a complete formalization (syntax and semantics) of SparkLite. In section 4 we describe the term semantics of SparkLite problem. In section 5.1 we provide a decision procedure for verifying equivalence of SparkLite programs without aggregate expressions. We continue in Section ??, where we discuss SparkLite programs with aggregate expressions: we present a sound equivalence verification technique as well as descriptions of classes of SparkLite programs for which the technique can be made complete.

2 Preliminaries

In this section, we describe a simple extension of Presburger arithmetic [15], which is the first-order theory of the natural numbers with addition, to tuples of integers, and state its decidability.

Notations. We denote the set of natural numbers (including zero), positive numbers, and integers by \mathbb{N} , \mathbb{N}^+ , and \mathbb{Z} , respectively. We write $ite(p, e, e')$ to denote an expression which evaluates to e if p holds and to e' otherwise. We use \perp to denote the *undefined* value. A *bag* m over a domain X is a multiset (i.e., a set which allows for repetitions) with elements taken from X . We write $\{\cdot\}$ and $\{\!\!\{\cdot\}\!\!\}$ to denote sets and bags, respectively. We denote the *size* (number of elements) of a set, respectively, a bag, X by $|X|$.

Presburger Arithmetic. We consider a fragment of first-order logic (FOL) with equality over the integers, where expressions are written in the syntax specified in Figure 1.¹ Disregarding the tuple expressions $((pe, \overline{pe})$ and $\mathbf{p}_i(e)$), the resulting first-order theory with the usual \forall and \exists quantifiers is called the *Presburger Arithmetic*.² The problem of checking whether a sentence in Presburger arithmetic is valid has long been known to be decidable [11, 15], even

¹ We assume the reader is familiar with FOL, and omit a more formal description for brevity.

² Originally, Presburger Arithmetic was defined as a theory over natural numbers. However, its extension to integers is also decidable. (See, e.g., [6].)

Arithmetic Exp.	$ae ::= c \mid ae + ae \mid -ae \mid c * ae \mid ae / c \mid ae \% c$
Boolean Exp.	$be ::= \mathbf{true} \mid \mathbf{false} \mid e = e \mid ae < ae \mid \neg be \mid be \wedge be \mid be \vee be$
Primitive Exp.	$pe ::= ae \mid be$
General Basic Exp.	$e ::= pe \mid v \mid (pe, \overline{pe}) \mid p_i(e) \mid \mathbf{ite}(be, e, e)$

■ **Figure 1** Terms of the Augmented Presburger Arithmetic

when combined with Boolean logic [6,13]. For example, *Cooper's Algorithm* [8] is a standard decision procedure for Presburger Arithmetic³.

In this paper, we consider a simple extension this language by adding a *tuple constructor* (pe, \overline{pe}) and a projection operator $p_i(e)$, and call the extended language *Augmented Presburger Arithmetic*. The decidability of Presburger Arithmetic, as well as Cooper's Algorithm, can be naturally extended to the Augmented Presburger Arithmetic. Intuitively, verifying the equivalence of tuple expressions can be done by verifying the equivalence their corresponding constituents.

► **Proposition 1.** *The theory of formulas over \mathbb{Z}^n with terms in the Augmented Presburger Arithmetic is decidable.*

3 The SparkLite language

In this section, we define the syntax of SparkLite, a simple imperative programming language which allows to use Spark's *resilient distributed datasets* (RDDs) [16].

3.1 Data Model

Basic types. SparkLite supports two primitive types: integers (`Int`) and booleans (`Boolean`). On top of this, the user can define types which are Cartesian products of primitive types. In the following we use c to range over integer numerals (constants), $b \in \{\mathbf{true}, \mathbf{false}\}$ to range over Boolean constants, and τ to range over basic types and record types.

RDDs. In addition, SparkLite allows the user to define *RDDs*. RDDs are bags of elements, all of the same type. Hence, RDD_τ denotes bags containing elements of type τ .

Semantic Domains. We interpret the integer and Boolean primitive types as *integers* (\mathbb{Z}) and *booleans* (\mathbb{B}), respectively. We use $\llbracket \cdot \rrbracket$ (semantic brackets) throughout the paper to denote the semantics of program constructs; for types, the semantics is a set of all the values pertaining to it. So $\llbracket \mathbf{Int} \rrbracket = \mathbb{Z}$, $\llbracket \mathbf{Boolean} \rrbracket = \mathbb{B}$.

The interpretation of both primitive types is denoted $T = \mathbb{Z} \cup \mathbb{B}$. The interpretation of all possible types (including mixed Cartesian products of the primitive types) is denoted by $\mathcal{T} = \bigcup_n T^n$.

An *RDD* type is interpreted as a *bag* (an unordered set allowing repeating elements). We write $\llbracket RDD_\tau \rrbracket = (\llbracket \tau \rrbracket \leftrightarrow \mathbb{N})$, meaning that an RDD value r of type RDD_τ is interpreted as a bag of values from τ , that is $\llbracket r \rrbracket \in (\llbracket \tau \rrbracket \leftrightarrow \mathbb{N})$. We let $\llbracket RDD \rrbracket = \bigcup_{\tau \in \mathcal{T}} \llbracket RDD_\tau \rrbracket$, the semantic domain of RDDs over all possible record types $\tau \in \mathcal{T}$.

³ The complexity of Cooper's algorithm is $O(2^{2^{pn}})$ for some $p > 0$ and where n is the number of symbols in the formula [14]. However, in practice, our experiments show that Cooper's algorithm on non-trivial formulas returns almost instantly, even on commodity hardware.

3.2 Functional Model

Operations. *RDDs* are analogous to database tables and as such the methods to query the *RDDs* are inspired by both *Relational Algebra (RA)* [4, 7] and Spark [17]. *RA* has 5 basic operators, which are *Select*, *Project*, *Cartesian Product*, *Union* and *Subtract*. This paper focuses on the first three operators. The *Select* operator is analogous to *filter* in SparkLite, and *Project* is analogous to *map*. The expressive power of SparkLite’s *map* and *filter* is greater than their analogous *RA* operations thanks to UDFs (see next), which allow *extended projection* as well as greater flexibility in executing complex operations on elements of different types.

UDFs. A special feature of Spark is allowing some of its standard operations to be *higher order functions* — they take a function and apply it to an *RDD* in a specific manner defined by the operation. For example, the operation *fold* can be applied to an *RDD* containing integer numbers by providing a function that adds to two integers to the *fold* transform, yielding the sum of all the numbers in the bag. Such a function is called a “*User-Defined Function*”, or *UDF*, for short. The *signature* of a *UDF* contains information on the return type and the arguments types, and when applied in the context of an *RDD* operation, the signature should match both the *RDD* type and the operation on it (see typing rules in appendix B). Each *UDF* has a *definition*, which takes the syntax $\lambda \bar{v}. e$, where \bar{v} are names of variables used as arguments. For example, $\text{addMod10} = \lambda x, y. x \% 10 + y \% 10$. This function’s *signature* would be $\text{addMod10}: \text{Int} \times \text{Int} \rightarrow \text{Int}$, that is, it takes two *Int* arguments and produces one *Int* value. The types are usually omitted from the definition for brevity, but when the types are not clear from the context, we may write them as annotations: $\text{addMod10} = \lambda x: \text{Int}, y: \text{Int}. x \% 10 + y \% 10: \text{Int}$. We could also write it as a function with single argument like this: $\text{addPairMod10} = \lambda z. p_1(z) \% 10 + p_2(z) \% 10$, where it would have the signature $\text{addPairMod10}: (\text{Int} \times \text{Int}) \rightarrow \text{Int}$. For readability, we sometimes give names to the projections and write them implicitly as $\text{addPairMod10} = \lambda(x, y). x \% 10 + y \% 10$; this is just a shortcut.

We allow the definition of these functions to be *parametric*, by using two levels of λ . The outer level denotes the parameters, which can be provided to obtain a regular function. For example, a function that adds 1 to an integer is written as $f = \lambda x. x + 1$, where a function that adds any constant is written as $g = \lambda a. \lambda x. x + a$. The function is *curried*, so that applying it to parameter values of the appropriate types produces a function (by *beta-reduction*): $g(1)$ is identical to f .

3.3 Syntax

The syntax of SparkLite language is defined in Figure 2.

Syntactic Categories. We assume variables to be an infinite syntactic category, ranged over by $v, b, r \in \text{Vars}$. Expressions range over e . An integer constant is denoted c . There are 4 operations: *map*, *filter*, *cartesian*, and *fold*. Some of the operations require arguments, which may be either a primitive expression, an *RDD*, or a function. Functions range over $f, F \in \text{LambdaExpressions}$. Parametric functions are denoted by capital meta-variables (F as opposed to f for regular functions) and must always be given the list of parameters when passed to an operation.

Basic Types	τ	$::=$	$\text{int} \mid \text{bool} \mid \tau \times \dots \times \tau$
RDDs	RDD	$::=$	RDD_τ
Variables	x	$::=$	$v \mid r$
Arithmetic Exp.	ae	$::=$	$c \mid ae + ae \mid -ae \mid c * ae \mid ae / c \mid ae \% c$
Boolean Exp.	be	$::=$	$\text{true} \mid \text{false} \mid e = e \mid ae < ae \mid \neg be \mid be \wedge be \mid be \vee be$
General Basic Exp.	e	$::=$	$ae \mid be \mid v \mid (e, e) \mid p_i(e) \mid \text{if } (b) \text{ then } e \text{ else } e$
Functions	$Fdef$	$::=$	$\text{def } f = \lambda \overline{y} : \overline{\tau} \ e : \tau$
Parametric Functions	$PFdef$	$::=$	$\text{def } F = \lambda \overline{x} : \overline{\tau}. \lambda \overline{y} : \overline{\tau} \ e : \tau$
RDD Exp.	re	$::=$	$\text{cartesian}(r, r) \mid \text{map}(f)(r) \mid \text{filter}(f)(r)$
RDD Aggregation Exp.	ge	$::=$	$\text{fold}(e, f)(r)$
General Exp.	η	$::=$	$e \mid re \mid ge$
Program Body	E	$::=$	$\text{Let } x = \eta \text{ in } E \mid \eta$
Program	$Prog$	$::=$	$P(\overline{r} : RDD_\tau, \overline{v} : \overline{\tau}) = \overline{Fdef} \ \overline{PFdef} \ E$

■ **Figure 2** Syntax for SparkLite

```

isOdd =  $\lambda x : \text{Int}. \neg(x \% 2 = 0)$ 
Let: doubleAndAdd =  $\lambda c : \text{Int}. \lambda x : \text{Int}. 2 * x + c$ 
    sumFlatPair =  $\lambda A : \text{Int}, (x, y) : \text{Int} \times \text{Int}. A + x + y$ 
1  A = filter(isOdd)(R0)
2  B = map(doubleAndAdd(1))(A)
3  C = cartesian(B, R1)
4  v = fold(0, sumFlatPair)(C)
5  return v

```

■ **Figure 3** Example SparkLite program

Program structure. The header of a program contains function definitions. Loops are not allowed in the body of a program. Variable declarations are in *SSA* (*Static Single Assignment*) form [9]. Variables are immutable by this construction. Programs have no side effects, do not change the inputs, and always return a value. The *program signature* will consist of its name, its input types and return type: $P(\overline{\tau}_i, RDD_i) : \tau_o$

Example program. Consider the example SparkLite program in Figure 3.

From the example program we can see the general structure of SparkLite programs: First, the functions that are used as UDFs in the program are declared and defined: *isOdd*, *sumFlatPair* defined as *Fdef*, and *doubleAndAdd* defined as a *PFdef*. The name of the program ($P = P1$) is announced with a list of input RDDs (R_0, R_1) (*Prog* rule). Instead of writing *Let l_1 in Let l_2 in ...*, we use syntactic sugar, where each line of code contains a single l_i , and the last line denotes the return value using the **return** keyword. Here, 3 variables of RDD type (A, B, C) and one integer variable (v) are bound by *Lets*. We can see in the definition of A an application of the *filter* operation, accepting the RDD R_0 and the function *isOdd*. For B 's definition we apply the *map* operation with a parametric function

$\llbracket c \rrbracket(\rho)$	$= c$
$\llbracket v \rrbracket(\rho)$	$= \rho(v)$
$\llbracket \text{unOp } e \rrbracket(\rho)$	$= \text{unOp } \llbracket e \rrbracket(\rho)$
$\llbracket e_1 \text{ binOp } e_2 \rrbracket(\rho)$	$= \llbracket e_1 \rrbracket(\rho) \text{ binOp } \llbracket e_2 \rrbracket(\rho)$
$\llbracket (e_1, \dots, e_n) \rrbracket(\rho)$	$= (\llbracket e_1 \rrbracket(\rho), \dots, \llbracket e_n \rrbracket(\rho))$
$\llbracket \text{if } e_1 \text{ then } e_2 \text{ else } e_3 \rrbracket(\rho)$	$= \text{ite}(\llbracket e_1 \rrbracket(\rho), \llbracket e_2 \rrbracket(\rho), \llbracket e_3 \rrbracket(\rho))$
$\llbracket \text{map}(f)(r) \rrbracket(\rho)$	$= \{\{\rho(f)(x) \mid x \in \rho(r)\}\}$
$\llbracket \text{filter}(b)(r) \rrbracket(\rho)$	$= \{\{x \mid x \in \rho(r) \wedge \rho(b)(x)\}\}$
$\llbracket \text{cartesian}(r_1, r_2) \rrbracket(\rho)$	$= \{\{(x_1, x_2) \mid x_1 \in \rho(r_1) \wedge x_2 \in \rho(r_2)\}\}$
$\llbracket \text{fold}(a_0, f)(r) \rrbracket(\rho)$	$= q(\llbracket a_0 \rrbracket(\rho), \rho(r)), \text{ where }$
	$q(v_0, s) = \begin{cases} v_0 & s = \emptyset \\ \rho(f)(x, q(v_0, s')) & s = \{x; 1\} \cup s' \end{cases}$
$\llbracket \text{Let } x = \eta \text{ in } E \rrbracket(\rho)$	$= \llbracket E \rrbracket(\rho[x \mapsto \llbracket \eta \rrbracket(\rho)])$
$\llbracket P(\dots) = \dots E \rrbracket(\rho_0)$	$= \llbracket E \rrbracket(\rho_0)$

■ **Figure 4** Semantics of SparkLite. $\text{Prog} = P(\overline{r}: \overline{RDD}_\tau, \overline{v}: \overline{\tau}) = \overline{Fdef} \ \overline{PFdef} \ E$. unOp and binOp are taken from Figure 2: $\text{unOp} \in \{-, \neg, \pi_i\}$, $\text{binOp} \in \{+, *, /, \%, =, <, \wedge, \vee, (,)\}$

doubleAndAdd with the parameter 1, which is interpreted as $\lambda x. 2 * x + 1$. C is the cartesian product of B and input RDD R_1 . We apply an aggregation using *fold* on the RDD C , with an initial value 0 and the function *sumFlatPair*, which ‘flattens’ elements of tuples in C , taking their sum. The sum total of all this elements is stored in the variable v . The returned value is the integer variable v . The program’s signature is $P1(RDD_{\text{Int}}, RDD_{\text{Int}}): \text{Int}$.

3.4 Operational Semantics

Program Environment. We define a unified semantic domain $\mathcal{D} = \mathcal{T} \cup RDD$ for all types in SparkLite. The *program environment* type:

$$\mathcal{E} = \text{Vars} \rightarrow \mathcal{D}$$

is a mapping from each variable in **Vars** to its value, according to type. A variable’s type does not change during the program’s run, nor does its value.

Data flow. We start with an initial environment function ρ_0 maps all input variables and function definitions. We define the *semantic interpretation* of expressions based on an environment $\rho \in \mathcal{E}$, and specifically for $x \in \overline{r} \cup \overline{v}$, $\llbracket x \rrbracket(\rho) = \rho(x)$. The semantics of composite expressions are straight-forward using the semantics of their components. The semantics of *Let* is to create a new environment by binding the variable name. In Figure 4 we specify the behavior of $\llbracket \cdot \rrbracket(\cdot)$ for all expressions and statements.

Function and UDF semantics. For UDFs, which are based on a restricted fragment of the *simply typed lambda calculus* $\llbracket \cdot \rrbracket$, we assume the syntax and semantics are the same as in the λ -calculus. Note however that the syntax does not allow passing higher order functions as UDFs, and forces any higher order function to be reduced to a first-order function beforehand. In addition, all parameters passed to UDF which are based on higher-order functions are read-only.

$$\begin{aligned}
\phi_P(\text{Let } x = \eta \text{ in } E, k) &= (t_{E'}[t_\eta/x], m), \text{ where } (t_{E'}, n) = \phi_P(E', k), (t_\eta, m) = \phi_P(\eta, n) \\
\phi_P(e, k) &= (e, k) \\
\phi_P(\text{map}(f)(r), k) &= (f(t), m), \text{ where } (t, m) = \phi_P(r, k) \\
\phi_P(\text{filter}(f)(r)) &= (\text{ite}(f(t) = tt, t, \perp), m), \text{ where } (t, m) = \phi_P(r, k) \\
\phi_P(\text{cartesian}(r_1, r_2), k) &= ((t_{r_1}, t_{r_2}), m), \text{ where } (t_{r_1}, n) = \phi_P(r_1, k), (t_{r_2}, m) = \phi_P(r_2, n) \\
\phi_P(\text{fold}(f, e)(r), k) &= ([t]_{e,f}, m), \text{ where } (t, m) = \phi_P(r, k) \\
\phi_P(r, k) &= \begin{cases} (\mathbf{x}_r^{(k)}, k+1) & r \in \bar{r} \\ (r, k) & \text{otherwise} \end{cases}
\end{aligned}$$

$$\begin{aligned}
\text{Let } P : P(\bar{r}, \bar{v}) &= \bar{F} \bar{f} E \\
\Phi(P) = t, & \text{ where } \phi_P(E, 0) = (t, _)
\end{aligned}$$

■ **Figure 5** Compiling SparkLite to logical terms (ϕ).

$$\begin{aligned}
\llbracket \mathbf{x}_{r_i}^{(k)} \rrbracket(\bar{r}) &= r_i \\
\llbracket f(t) \rrbracket(\bar{r}) &= \{ \llbracket f \rrbracket(z) \mid z \in \llbracket t \rrbracket(\bar{r}) \} \\
\llbracket \text{ite}(f(t), t, \perp) \rrbracket(\bar{r}) &= \{ z \mid z \in \llbracket t \rrbracket(\bar{r}) \wedge \llbracket f \rrbracket(z) \} \\
\llbracket (t, t') \rrbracket(\bar{r}) &= \{ (z, z') \mid z \in \llbracket t \rrbracket(\bar{r}) \wedge z' \in \llbracket t' \rrbracket(\bar{r}) \} \\
\llbracket [t]_{e,f} \rrbracket(\bar{r}) &= \llbracket \text{fold}(e, f)(r') \rrbracket[r' \mapsto \llbracket t \rrbracket(\bar{r})]
\end{aligned}$$

■ **Figure 6** Semantics of terms.

Semantics of operations. In Figure 4 the semantics of all RDD expressions, including aggregation, are explicitly stated.

Notes. In the *bag semantics*:

- In *map*, if f maps y to x , the multiplicity of x is the sum of multiplicities of all y elements. In other words, if an element x appears n times, we apply the f on it n times.
- *fold* is *well defined*: A fold UDF f should satisfy $\forall A, x, y. f(f(A, x), y) = f(f(A, y), x)$ (order of elements does not affect the result).

Example. For the example program Figure 3, suppose we were given the following input: $R_0 = \{(1; 7), (2; 1)\}$, $R_1 = \{(3; 4), (5; 2)\}$. Then: $\rho(A) = \{(1; 7)\}$, $\rho(B) = \{(3; 7)\}$, $\rho(C) = \{((3, 3); 28), ((3, 5); 14)\}$, $\rho(v) = 28 * (3+3) + 14 * (3+5)$ and the program returns $\rho(v) = 280$.

4 Term Semantics for SparkLite

In this section, we present an alternative, equivalent semantics for SparkLite where the program is interpreted as a term in the Augmented Presburger Arithmetic. This term is called the *program term* and denoted $\Phi(P)$ for program P , specified in Figure 5. The variables of the term are taken from the input RDDs. We take the previously analyzed example program from Figure 3. The Φ function is defined using the function ϕ whose purpose is to maintain unique variable names. ϕ is applied recursively on the expression

XX:8 Verifying Equivalence of Spark Programs

returned by the program. We simplify by running ϕ_{P1} on each line of the program, top-down:

$$\begin{aligned}
\phi_P(A, 0) &= \phi_P(\text{filter}(\text{isOdd})(R_0), 0) = (\text{ite}(\text{isOdd}(t), t, \perp), n \text{ where } (t, n) = \phi_P(R_0, 0)) \\
&= \phi_P(R_0, 0) = (\mathbf{x}_{R_0}^{(0)}, 1) \quad (\text{ite}(\text{isOdd}(\mathbf{x}_{R_0}^{(0)}), \mathbf{x}_{R_0}^{(0)}, \perp), 1) \\
\phi_P(B, 1) &= (\phi_P(\text{doubleAndAdd}(1)(A), 1) = \text{doubleAndAdd}(1)(A), 1) \\
&= (\text{doubleAndAdd}(1)(\text{ite}(\text{isOdd}(\mathbf{x}_{R_0}^{(0)}), \mathbf{x}_{R_0}^{(0)}, \perp)), 1) \\
\phi_P(C, 1) &= \phi_P(\text{cartesian}(B, R_1), 1) = \phi_P(B, 1) = (B, 1) = ((B, p_1(\phi_P(R_1, 1))), p_2(\phi_P(R_1, 1))) \\
&= \phi_P(R_1, 1) = (\mathbf{x}_{R_1}^{(1)}, 2) \quad ((B, \mathbf{x}_{R_1}^{(1)}), 2) \\
\phi_P(v, 2) &= \phi_P(\text{fold}(0, \text{sumFlatPair})(C), 2) = \phi_P(C, 2) = (C, 2) \quad ([C]_{0, \text{sumFlatPair}}, 2) \\
\Phi(P) &= p_1(\phi_P(v, 2)) = [C]_{0, \text{sumFlatPair}} = [(B, \mathbf{x}_{R_1}^{(1)})]_{0, \text{sumFlatPair}} \\
&= [(\text{doubleAndAdd}(1)(\text{ite}(\text{isOdd}(\mathbf{x}_{R_0}^{(0)}), \mathbf{x}_{R_0}^{(0)}, \perp)), \mathbf{x}_{R_1}^{(1)})]_{0, \text{sumFlatPair}}
\end{aligned}$$

Representative elements of RDDs. The variables assigned by ϕ for input RDDs are called *representative elements*. In a program that receives an input RDD r , we denote the representative element of r as: \mathbf{x}_r . The set of possible valuations to that variable is equal to the bag defined by r , and an additional ‘undefined’ value (\perp), for the empty RDD. Therefore \mathbf{x}_r ranges over $\text{dom}(r) \cup \{\perp\}$. By abuse of notation, the term for a non-input RDD, computed in a SparkLite program, is also called a representative element.

Comparing representative elements. For two program terms to be comparable, they must depend on the same input RDDs. For example:

	$P1(R_0: RDD_{\text{Int}}, R_1: RDD_{\text{Int}}):$	$P2(R_0: RDD_{\text{Int}}, R_1: RDD_{\text{Int}}):$
1	<code>return map($\lambda x.1$)(R_0)</code>	<code>return map($\lambda x.1$)(R_1)</code>

$P1$ and $P2$ have the same program term (the constant 1), but the multiplicity of that element in the output bag is different and depends on the source input RDD. In $P1$, its multiplicity is the same as the size of R_0 , and in $P2$ it is the same as the size of R_1 . $P1$ and $P2$ are therefore not equivalent, because we can provide inputs R_0, R_1 of different sizes. Therefore, for each program term $\Phi(P)$ we consider the *set of free variables*, $FV(\Phi(P))$. Each free variable has some source input RDD, and an input RDD may have more than one free variable representing it in the program term. In the example, $FV(\Phi(P1)) = \{\mathbf{x}_{R_0}\}$, and $FV(\Phi(P2)) = \{\mathbf{x}_{R_1}\}$. For programs to be equivalent, there must be an *isomorphism* between the sets, mapping each free variable to a single free variable with the same source input RDD. An exception to the rule of having an isomorphism of the free variables are *trivial programs*, that always return the empty RDD, which has no multiplicity.

Formalization of the Term Semantics for SparkLite. Let P be a SparkLite program. We use standard notations \bar{r} for the inputs of P , and r^{out} for the output of P . The term $\Phi(P)$ is called the *program term* of P as before. We write the set of free variables of the program term $FV(\Phi(P))$ as a vector: $(\mathbf{x}_{r_{j_k}})_{k=1}^{n_P}$, where n_P is the number of free variables. A vector of valuations to the free variables is denoted $\bar{x} = (x_1, \dots, x_{n_P})$ and satisfies $x_k \in r_{j_k}$ for $k \in \{1, \dots, n_P\}$. The *Term Semantics* (TS) of a program that returns an RDD-type output is the bag that is obtained from all possible valuations to the free variables:

$$TS(P)(\bar{r}) = \{ \{ \Phi(P)[\bar{x}/FV(\Phi(P))] \mid \Phi(P)[\bar{x}/FV(\Phi(P))] \neq \perp \wedge \forall k \in \{1, \dots, n_P\}. x_k \in r_{j_k} \} \}$$

Assigning a concrete valuation to the free variables of $\Phi(P)$ returns an element in the output RDD r^{out} . By taking all possible valuations to the term with elements from \bar{r} , we get the bag equal to r^{out} . For a program that returns a basic type and not an RDD, $TS(P) = \Phi(P)$.

► **Proposition 2.** Let $P : P(\bar{r}) = \overline{F} \overline{f} E$ be a SparkLite program, $\llbracket P \rrbracket$ be the interpretation of its output according to the operational semantics, and $TS(P)$ by the term semantics of P . Then, for any input \bar{v}, \bar{r} , we have:

$$TS(P)(\bar{v}, \bar{r}) = \llbracket P \rrbracket(\llbracket \bar{v} \rrbracket, \llbracket \bar{r} \rrbracket)$$

5 Verifying Equivalence of SparkLite Programs

The Program Equivalence (PE) problem. Let P_1 and P_2 be SparkLite programs, with signature $P_i(\bar{T}, RDD_T) : \tau$ for $i \in \{1, 2\}$. We use $\llbracket P_i \rrbracket(\llbracket \bar{v} \rrbracket, \llbracket \bar{r} \rrbracket)$ to denote the result of P_i . We say that P_1 and P_2 are *equivalent*, if for all input values \bar{v} and RDDs \bar{r} , it holds that $\llbracket P_1 \rrbracket(\llbracket \bar{v} \rrbracket, \llbracket \bar{r} \rrbracket) = \llbracket P_2 \rrbracket(\llbracket \bar{v} \rrbracket, \llbracket \bar{r} \rrbracket)$.

5.1 Verifying Equivalence of SparkLite Programs without Aggregations

Program equivalence problem formalization in TS semantics. Given two programs P, Q receiving as input a series of RDDs $\bar{r} = (r_1, \dots, r_n)$. We assume w.l.o.g. the programs are non-trivial, meaning they do not return the empty RDD for any choice of inputs. We define isomorphism of sets of free variables for P, Q as an injective and onto mapping: $S : FV(\Phi(P)) \xrightarrow{\sim} FV(\Phi(Q))$ such that $\forall k \in \{1, \dots, n\}. S(\mathbf{x}_{r_{jk}}^{(k)}) = \mathbf{x}_{r_{ji}}^{(i)} \wedge j_k = j_i$

The *PE* problem becomes the problem of proving the following:

$$\begin{aligned} (*) \quad & FV(\Phi(P)) \simeq^S FV(\Phi(Q)) \\ (**) \quad & \forall \bar{x}. \Phi(P)[\bar{x}/FV(\Phi(P))] = \Phi(Q)[\bar{x}/S(FV(\Phi(P)))] \end{aligned}$$

where the choice of \bar{r} is arbitrary, and S is non-deterministically chosen from all legal isomorphisms.

► **Lemma 1** (Decidability for programs without aggregations). *Given two SparkLite programs P and Q which do not contain aggregate operations, PE is decidable.*

Proof. For non-RDD return types, the absence of aggregate operators implies we can use Proposition 1, as the returned expression is expressible in the Augmented Presburger Arithmetic. For RDDs we provide an algorithm in Figure 7, which is a decision procedure. The correctness of the algorithm follows from the equivalence of the TS semantics and the operational semantics defined in 3.4. The algorithm generates an equivalence formula from the program terms of P and Q , which is a formula in the Augmented Presburger Arithmetic. Thus, its decidability again follows from Proposition 1. ◀

Examples. Note: All examples use syntactic sugar for ‘Let’ expressions. For brevity, instead of applying ϕ on the underlying ‘Let’ expressions, we apply it line-by-line from the top-down. Finding the isomorphism S between variable names is done automatically in all examples, but formally it is part of the decision procedure. In addition, we assume that in programs returning an RDD-type, the RDD is named r^{out} , and the programs always end with **return** r^{out} . Thus, $\Phi(P) = p_1(\phi_P(r^{out}))$.

► **Example 1** (Basic optimization - operator pushback). This example shows a common optimization of pushing the filter/selection operator backward, to decrease the size of the dataset.

	$P1(R: RDD_{Int}):$	$P2(R: RDD_{Int}):$
1	$R' = \text{map}(\lambda x. 2 * x)(R)$	$R' = \text{filter}(\lambda x. x < 7)(R)$
2	return $\text{filter}(\lambda x. x < 14)(R')$	return $\text{map}(\lambda x. x + x)(R')$

XX:10 Verifying Equivalence of Spark Programs

1. If $\Phi(P) = \emptyset \wedge \Phi(Q) = \emptyset$, output **equivalent**.
2. Verify that:

$$FV(\Phi(P)) = FV(\Phi(Q))$$

If not, output **not equivalent**.

3. a. Choose an isomorphism \mathcal{S} of the representative elements of the input RDDs in both P, Q .
- b. We check the following formula is satisfiable:

$$\exists \bar{v}. \Phi(P)[\bar{v}/FV(\Phi(P))] \neq \Phi(Q)[\bar{v}/S(FV(\Phi(Q)))]$$

- c. If it is satisfiable, go back to (a) and repeat until finding an unsatisfiable formula, or all possible isomorphisms were exhausted.
- d. If the formula is unsatisfiable, return **equivalent**.
- e. If all isomorphisms were exhausted without finding an unsatisfiable formula, then return **not equivalent**.

■ **Figure 7** An algorithm for solving PE for two programs P, Q with the same signature

Non trivial programs. Both programs may return an non-empty RDD of integers.

Free variables: $FV(\Phi(P1)) = \{\mathbf{x}_R\} = FV(\Phi(P2))$. The sets of free variables are equal.

Analysis of representative elements:

$$\begin{aligned} \phi_{P1}(R') &= 2 * \mathbf{x}_R; & \phi_{P1}(r^{out}) &= ite(\varphi < 14 \wedge \varphi = \phi_{P1}(R'), \varphi, \perp) = ite(2 * \mathbf{x}_R < 14, 2 * \mathbf{x}_R, \perp) \\ \phi_{P1}(R') &= ite(\mathbf{x}_R < 7, \mathbf{x}_R, \perp); & \phi_{P2}(r^{out}) &= (\lambda x. x + x)(\phi_{P1}(R')) = ite(\mathbf{x}_R < 7, \mathbf{x}_R, \perp) + ite(\mathbf{x}_R < 7, \mathbf{x}_R, \perp) \end{aligned}$$

We need to verify that:

$$\forall \mathbf{x}_R. ite(2 * \mathbf{x}_R < 14, 2 * \mathbf{x}_R, \perp) = ite(\mathbf{x}_R < 7, \mathbf{x}_R, \perp) + ite(\mathbf{x}_R < 7, \mathbf{x}_R, \perp)$$

To prove this, we need to encode the cased expressions in Presburger arithmetic. Undefined (\perp) values indicate ‘don’t care’ and are not part of the Presburger arithmetic. However, they can be handled by assuming the ‘if’ condition is satisfied, and verifying that the condition is indeed satisfied equally for all inputs. The first condition, therefore, is that both ‘if’ conditions agree on all possible values. The second condition is that when the condition holds, the resulting expressions are equivalent.

► **Proposition 3 (Schemes for converting conditionals to a normal form).** *We write a series of universally true schemes for translating cased expressions to Presburger arithmetic when appearing in an equivalence formula:*

1. The following useful identity for applying functions on a conditional is true:

$$f(ite(cond, e, \perp)) = ite(cond, f(e), \perp)$$

2. Equivalence of functions of conditionals:

$$(f(ite(cond, e, \perp)) = g(ite(cond', e', \perp))) \iff ((cond \iff cond') \wedge (cond \implies f(e) = g(e')))$$

3. *Equivalence of a function of a conditional and an arbitrary expression:*

$$(f(ite(cond, e, \perp)) = e') \iff (cond \wedge f(e) = e')$$

4. *Applying a function with multiple arguments on multiple conditionals (a function receiving \perp input as one of its arguments returns a \perp):*

$$f(ite(cond, e, \perp), ite(cond', e', \perp)) = ite(cond \wedge cond', f(e, e'), \perp)$$

5. *Applying a function with multiple arguments on a conditional and a general expression:*

$$f(ite(cond, e, \perp), e') = ite(cond, f(e, e'), \perp)$$

The two last rules define the base case for functions with more than 2 arguments where at least one of the arguments is a conditional.

6. *Unnesting of nested conditionals*

$$ite(c_{ext}, ite(c_{int}, e, \perp), \perp) = ite(c_{int} \wedge c_{ext}, e, \perp)$$

Using Cooper's algorithm and the above schemes, we can prove the equivalence formula is true. See ?? for an implementation. ■

For additional examples, refer to appendix ??.

5.2 Verifying Equivalence of a Class of SparkLite Programs with Aggregation

In the following section we discuss how the existing framework can be extended to prove equivalence of SparkLite programs containing aggregate expressions. As mentioned, for programs returning a non-RDD value, $TS(P) = \Phi(P)$. In particular, it may do so by applying the *fold* operation on one or more RDDs. The terms for aggregate operations are given using a special operator - $[t]_{i,f}$ - where t is the term being folded, i is the initial value, and f is the fold function. The operator *binds* all free variables in the term t .

5.2.1 Single aggregate

The simplest class of programs in which an aggregation operator appears, is where a single aggregate operation is performed and it is the last RDD operation.

► **Example 2 (Maximum and minimum).** Below is an example of two equivalent programs representing the simplest class of programs with aggregation:

Let: $max = \lambda M, x. \text{if}(x > M) \text{ then } \{x\} \text{ else } \{M\}$ $min = \lambda M, x. \text{if}(x < M) \text{ then } \{x\} \text{ else } \{M\}$		
1	P1($R : RDD_{Int}$): return $\text{fold}(\perp, max)(R)$	P2($R : RDD_{Int}$): $R' = \text{map}(\lambda x. -x)(R)$
2		return $\text{fold}(\perp, min)(R')$

The programs compute the maximum element of a numeric RDD in two different methods: in the first program by getting the maximum directly, and in the second by getting the additive inverse of the minimum of the additive inverses of the elements. The equivalence formula is:

$$[\mathbf{x}_R]_{\perp, max} = -[-\mathbf{x}_R]_{\perp, min}$$

XX:12 Verifying Equivalence of Spark Programs

The two program apply a fold operation on two RDDs of equal size. We use an inductive claim to prove the equivalence:

$$\forall x, A, A'. A = -A' \Rightarrow \max(A, x) = -\min(A', -x)$$

The inductive claim is a formula in the Augmented Presburger Arithmetic. We assume $A = -A'$ and attempt to prove:

$$\begin{aligned} \max(A, x) & \stackrel{?}{=} -\min(A', -x) \\ \begin{cases} A & A > x \\ x & \text{otherwise} \end{cases} & \stackrel{?}{=} \begin{cases} A' & A' < -x \\ -x & \text{otherwise} \end{cases} = \begin{cases} -A' & A' < -x \\ x & \text{otherwise} \end{cases} \stackrel{A=-A'}{=} \begin{cases} A & -A < -x \\ x & \text{otherwise} \end{cases} = \begin{cases} A & A > x \\ x & \text{otherwise} \end{cases} \end{aligned}$$

Indeed, by replacing $A' = -A$ we get equal expressions. ■

5.2.1.1 Soundness

We present a generalization of the inductive claim for the class of programs discussed here, of programs performing a single aggregate operation after a series of *map*, *filter* and *cartesian* operations. Those definitions lead to the following lemma:

► **Lemma 2** (Sound method for verifying equivalence of aggregate terms). *Let there be representative elements φ_0, φ_1 over σ_0, σ_1 . We assume φ_0, φ_1 were composed from *map*, *filter* and *cartesian product* (without self products⁴). Let there be two fold functions $f_0 : \xi_0 \times \sigma_0 \rightarrow \xi_0, f_1 : \xi_1 \times \sigma_1 \rightarrow \xi_1$, two initial values $init_0 : \xi_0, init_1 : \xi_1$, and two functions $g : \xi_0 \rightarrow \xi, g' : \xi_1 \rightarrow \xi$. We have: if*

$$FV(\varphi_0) \simeq FV(\varphi_1), \text{ denoted } FV \tag{1}$$

$$g(init_0) = g'(init_1) \tag{2}$$

$$\begin{aligned} \forall \bar{v}, A_{\varphi_0} : \xi_0, A_{\varphi_1} : \xi_1. g(A_{\varphi_0}) = g'(A_{\varphi_1}) \implies \\ g(f_0(A_{\varphi_0}, \varphi_0[\bar{v}/FV])) = g'(f_1(A_{\varphi_1}, \varphi_1[\bar{v}/FV])) \end{aligned} \tag{3}$$

$$\text{then } g([\varphi_0]_{init_0, f_0}) = g'([\varphi_1]_{init_1, f_1})$$

Lemma 2 shows that an inductive proof of the equality of folded values is *sound*. Therefore, given two folded expressions which are not equivalent, the lemma is guaranteed to report so.

Additional Examples. Refer to appendix ??.

5.2.1.2 Completeness

There are several cases in which one or more of the requirements of Lemma 2 are not satisfied, yet the aggregate expressions are equal. The first requirement, $FV(\varphi_0) \simeq FV(\varphi_1)$, is not necessary when the fold functions applied on the terms are trivial. Suppose:

$$\forall \bar{v}. f_0(init_0, \varphi_0[\bar{v}/FV(\varphi_0)]) = init_0 \wedge \forall \bar{v}'. f_1(init_1, \varphi_1[\bar{v}'/FV(\varphi_1)]) = init_1$$

⁴ Self-cartesian products require multiple induction steps. For example, for $R \times R$, the existence of an element (x, y) such that $x \neq y$, implies the existence of an element (y, x) . Thus, the number of steps applied each time is determined according to the number of symmetric elements of a certain valuation. Due to lack of space, we omit this discussion from the paper.

Then:

$$[\varphi_0]_{init_0, f_0} = init_0 \wedge [\varphi_1]_{init_1, f_1} = init_1$$

And by Equation (3), equivalence follows. Conversely, when the fold is not trivial, the proof of lemma 2 requires the sets of free variables to be isomorphic. Otherwise, the induction termination is not well defined. One possibility is that the size of participating RDDs may not be equal. Assuming one set of free variables may be embedded in the other (the embedding is an isomorphism on the smaller set), any non-constant result of the fold function on these additional elements will lead to unequal results. To avoid such peculiarities, we shall consider only classes of programs that satisfy Equation (1). We proceed with an example showing that when we apply on folded expressions a non-injective function in both programs, the lemma may fail:

► **Example 3 (Non-injective modification of folded expressions).** Non-injective transformations can weaken the inductive claim, resulting in failure to prove it. As a result, lemma 2 fails to prove the equivalence of the following two programs.

	P1($R: RDD_{Int}$):	P2($R: RDD_{Int}$):
1	$R' = \text{map}(\lambda x.x\%3)(R)$	$R' = \text{fold}(0, \lambda A, x.A + x)(R)$
2	$\text{return fold}(0, \lambda A, x.(A + x)\%3)(R') = 0$	$\text{return } R'\%3 = 0$

To prove the equivalence, we should check by induction the equality of both boolean results. Taking $g(x) = \lambda x.x = 0$, $g'(x) = \lambda x.(x \bmod 3) = 0$ and attempt to prove by induction the following claim:

$$[x \bmod 3]_{0, + \bmod 3} = 0 \Leftrightarrow [x]_{0, + \bmod 3 \bmod 3} = 0$$

$$\forall x, A, A'. A = 0 \Leftrightarrow A' \bmod 3 = 0 \implies (A + x \bmod 3) \bmod 3 = 0 \Leftrightarrow (A' + x) \bmod 3 = 0$$

fails, the counter-example being $A = 1, A' = 2$. Then the hypothesis that says $A = 0 \Leftrightarrow A' \bmod 3 = 0$ is satisfied, but it cannot be said that $(A + x \bmod 3) \bmod 3 = 0 \Leftrightarrow (A' + x) \bmod 3 = 0$ (take $x = 1$: $((1 + (1\%3))\%3 = 2, (2 + 1)\%3 = 0)$).

From the above example we derive the first completeness criterion:

► **Lemma 3 (A decidable class of programs with a single aggregate expression).** *Under the premises of Lemma 2, we look at the class of SparkLite programs such that:*

$$FV(\varphi_0) \simeq FV(\varphi_1), \text{ denoted } FV \quad (1)$$

$$\begin{aligned} \forall \bar{v}_1, \bar{v}_2. \exists \bar{v}'. f_0(f_0(init_0, \varphi_0[\bar{v}_1/FV]), \varphi_0[\bar{v}_2/FV]) &= f_0(init_0, \varphi_0[\bar{v}'/FV]) \\ \wedge f_1(f_1(init_1, \varphi_1[\bar{v}_1/FV]), \varphi_1[\bar{v}_2/FV]) &= f_1(init_1, \varphi_1[\bar{v}'/FV]) \end{aligned} \quad (2)$$

For these programs, $g([\varphi_0]_{init_0, f_0}) = g'([\varphi_1]_{init_1, f_1})$ if and only if:

$$g(init_0) = g'(init_1) \quad (3)$$

$$\begin{aligned} \forall v, y. (A_{\varphi_0} = f_0(init_0, \varphi_0[y/FV]) \wedge A_{\varphi_1} = f_1(init_1, \varphi_1[y/FV]) \wedge g(A_{\varphi_0}) = g'(A_{\varphi_1})) \\ \implies g(f_0(A_{\varphi_0}, \varphi_0[v/FV])) = g'(f_1(A_{\varphi_1}, \varphi_1[v/FV])) \end{aligned} \quad (4)$$

Proof. Sound (if): We prove the equality $g([\varphi_0]_{init_0, f_0}) = g'([\varphi_1]_{init_1, f_1})$ by induction on the number of possible valuations v to FV , denoted n . For $n = 0$, $R_0 = R_1 = \emptyset$, thus $[\varphi_i]_{init_i, f_i} = init_i$ ($i = 1, 2$), and the equality follows from Equation (3). Assuming for n and proving for $n+1$: We let a sequence of intermediate values $A_{\varphi_i, k}$, ($i = 1, 2; k = 1, \dots, n+1$), for which we know in particular that $g(A_{\varphi_0, n}) = g'(A_{\varphi_1, n})$, and we need to prove $g(A_{\varphi_0, n+1}) = g'(A_{\varphi_1, n+1})$. We denote $A_{\varphi_i, 0} = init_i$, and then we have $A_{\varphi_i, k} = f_i(A_{\varphi_i, k-1}, a_k)$ ($k = 1, \dots, n+1$).

1) for some a_i . According to Equation (2), $A_{\varphi_i,2} = f_i(A_{\varphi_i,1}, a_2) = f_i(f_i(\text{init}_i, a_1), a_2)$ yields $\exists a'_2. \bigwedge_{i=1,2} A_{\varphi_i,2} = f_i(\text{init}_i, a'_2)$. We can thus use Equation (2) to prove by induction that $\exists a'_k. \bigwedge_{i=1,2} A_{\varphi_i,k} = f_i(\text{init}_i, a'_k)$, and in particular $\exists a'_n. \bigwedge_{i=1,2} A_{\varphi_i,n} = f_i(\text{init}_i, a'_n)$. By applying Equation (4) for $v = a_{n+1}, y = a'_n$, we get:

$$\begin{aligned} g(f_0(f_0(\text{init}_0, \varphi_0[y/FV], \varphi_0[v/FV])) &= g'(f_1(f_1(\text{init}_1, \varphi_1[y/FV], \varphi_1[v/FV])) &\implies \\ g(f_0(f_0(\text{init}_0, \varphi_0[a'_n/FV], \varphi_0[v/FV])) &= g'(f_1(f_1(\text{init}_1, \varphi_1[a'_n/FV], \varphi_1[v/FV])) &\implies \\ g(f_0(A_{\varphi_0,n}, \varphi_0[a_{n+1}/FV])) &= g'(f_1(A_{\varphi_1,n}, \varphi_1[a_{n+1}/FV])) &\implies \\ g(A_{\varphi_0,n+1}) &= g'(A_{\varphi_1,n+1}) \end{aligned}$$

as required.

Complete (only if): Assume towards a contradiction that either Equations (3) and (4) are false. If the requirement of Equation (3) is not satisfied, yet the aggregates are equivalent, i.e.

$$g([\varphi_0]_{\text{init}_0, f_0}) = g'([\varphi_1]_{\text{init}_1, f_1}) \wedge g(\text{init}_0) \neq g'(\text{init}_1)$$

then we can get a contradiction by choosing all input RDDs to be empty. Thus, $[\varphi_0]_{\text{init}_0, f_0} = \text{init}_0 \wedge [\varphi_1]_{\text{init}_1, f_1} = \text{init}_1 \implies g(\text{init}_0) = g'(\text{init}_1)$, contradiction. The conclusion is that Equation (3) is a necessary condition for equivalence. Therefore, we assume just Equation (4) is false. Let there be counter-examples \bar{v}, \bar{y} to Equation (4), and let:

$$F_i = f_i(f_i(\text{init}_i, \varphi_i[\bar{y}/FV], \varphi_i[\bar{v}/FV]))$$

Then $g(F_0) \neq g'(F_1)$. By Equation (2) we can write F_i as: $F_i = f_i(\text{init}_i, \varphi_i[\bar{w}/FV])$ for some \bar{w} . We take an RDD $R = \{\{\bar{w}\}\}$ with multiplicity 1. Then $\llbracket \varphi_j \rrbracket(R) = \{\{\varphi_j(\bar{w})\}\}$, for which: $\llbracket [\varphi_j]_{\text{init}_j, f_j} \rrbracket(R) = F_i$. By the assumption, $g(\llbracket [\varphi_j]_{\text{init}_j, f_j} \rrbracket(R)) = g'(\llbracket [\varphi_j]_{\text{init}_j, f_j} \rrbracket(R))$, but then $g(F_0) = g'(F_1)$. Contradiction. \blacktriangleleft

► **Example 4 (Completing Example 3).** We have:

$$\begin{aligned} f_0(f_0(0, x \% 3), y \% 3) &= x \% 3 + y \% 3 = f_0(0, (x + y) \% 3) = (x + y) \% 3 \% 3 \\ f_1(f_1(0, x), y) &= x + y = f_1(0, x + y) \end{aligned}$$

So Equation (2) is true. Indeed, we use the same v , $v = x + y$, to reduce the nesting of the fold functions for arbitrary x, y . We are left with proving:

$$\forall x, y. ((0 + y \% 3) \% 3 = 0 \iff (0 + y) \% 3 = 0) \implies (y + x \% 3) \% 3 = 0 \iff (y + x) \% 3 = 0$$

which is correct — so we were able to prove the equivalence with the stronger lemma.

► **Lemma 4 (Lifting to programs with terms dependent on an aggregation).** *Let there be two SparkLite programs P_1, P_2 containing two aggregated expressions $a_i = [\varphi_i]_{\text{init}_i, f_i}$ for $i \in \{1, 2\}$. Let the program terms of P_1, P_2 be $\psi_j(a_j)$ with free variables $FV(\psi_j)$. P_1 is equivalent to P_2 if:*

$$FV(\varphi_1) \simeq FV(\varphi_2), \text{ denoted } FV_\varphi \tag{1}$$

$$FV(\psi_1) \simeq FV(\psi_2), \text{ denoted } FV_\psi \tag{2}$$

$$\forall \bar{x}. \psi_1(\text{init}_1)[\bar{x}/FV_\psi] = \psi_2(\text{init}_2)[\bar{x}/FV_\psi] \tag{3}$$

$$\forall \bar{v}, A_1, A_2. (\forall \bar{x}. \psi_1(A_1)[\bar{x}/FV_\psi] = \psi_2(A_2)[\bar{x}/FV_\psi]) \implies \tag{4}$$

$$(\forall \bar{x}. \psi_1(f_1(A_1, \varphi_1[\bar{v}/FV_\varphi]))[\bar{x}/FV_\psi] = \psi_2(f_2(A_2, \varphi_2[\bar{v}/FV_\varphi]))[\bar{x}/FV_\psi]) \tag{5}$$

Lemmas 3,4 show classes of programs in which the inductive argument can be used to prove equivalence of the aggregate expressions.⁵

5.2.2 A class for which *PE* is undecidable

We show a reduction of Hilbert's 10th problem to *PE*. We assume towards a contradiction that *PE* is decidable under the premises of lemma 5 with representative elements based also on self-products. Let there be a polynomial p over k variables x_1, \dots, x_k , and coefficients a_1, \dots, a_k . For each variable x_i we assume the existence of some RDD R_i with x_i elements. We use SparkLite operations and the input RDDs R_i to represent the value of the polynomial P for some valuation of the x_i . For each summand in the polynomial p , we define a translation φ :

must be after multi independent folds (lemma 7) which is referenced here!

- $x_i \longrightarrow^\varphi [\text{map}(\lambda x.1)(R_i)]_{0, \lambda A, x. A+x}$
- $x_i x_j \longrightarrow^\varphi [\text{cartesian}(\text{map}(\lambda x.1)(R_i), \text{map}(\lambda x.1)(R_j))]_{0, \lambda A, (x,y). A+x}$
- By induction, $x_i^2 \longrightarrow^\varphi [\text{cartesian}(\text{map}(\lambda x.1)(R_i), \text{map}(\lambda x.1)(R_i))]_{0, \lambda A, (x,y). A+x}$.
This rule as well as the rest of the powers follow according to the previous rule. For a degree k monom, we apply the *cartesian* operation k times, and fold it with $\lambda A, (x)_{i=1}^k. A + x_1$.
- $x_i^0 \longrightarrow^\varphi 1$, trivially.
- $am(x) \longrightarrow^\varphi a\varphi(m(x))$ where $m(x)$ is a monomial with coefficient 1 of the variable x , thus we have already defined φ for it.
- $aq(x_{i_1}, \dots, x_{i_j}) \longrightarrow^\varphi a\varphi(q(x_{i_1}, \dots, x_{i_j}))$ where $q(x)$ is a monomial with coefficient 1 and multiple variables for which φ was defined in the previous rules.
- $\varphi(p(a_1, \dots, a_k; x_1, \dots, x_k)) = \sum_{i=1}^k \varphi(a_i q(x_{i_1}, \dots, x_{i_k}))$ follows by structural induction on the previous rules.

We generate the following instance of the *PE* problem:

	$P1(R_1, \dots, R_k: RDD_{\text{Int}}):$	$P2(R_1, \dots, R_k: RDD_{\text{Int}}):$
1	return $\varphi(p) \neq 0$	return tt

By choosing input RDDs such that the size of R_i is equal to the matching variable x_i , we can simulate any valuation to the polynomial p . If $P1$ returns true, then the valuation is not a root of the polynomial p . Thus, if it is equivalent to the 'true program' $P2$, then the polynomial p has no roots. Therefore, if the algorithm solving *PE* outputs 'equivalent' then the polynomial p has no root, and if it outputs 'not equivalent' then the polynomial p has some root, where $x_i = ||R_i||$. Thus we have polynomial reduction to Hilbert's 10th problem.

5.2.3 Multiple aggregates

5.2.3.1 Independent multiple aggregations.

Another relatively simple case of SparkLite programs is when the program contains multiple aggregate operations, but they are independent of each other. Namely, the result of one aggregate operations is not used in the other one.

⁵ Even when none of the completeness criteria are met, we may be successful in proving the equivalence using Lemma 2. For example, $[((\lambda x.1)(\mathbf{x}_{r_0}), (\lambda x.1)(\mathbf{x}_{r_1}))]_{0, \lambda A, (x,y). A+x+y} = [((\lambda x.1)(\mathbf{x}_{r_1}), (\lambda x.1)(\mathbf{x}_{r_0}))]_{0, \lambda A, (x,y). A+x+y}$ can be proved by induction, but for $f = \lambda A, (x, y). A+x+y$, $f(f(A, (1, 1)), (1, 1)) = A+4 \neq f(A, (1, 1)) = A+2$ (the choice of valuation does not change the result). Thus, it does not satisfy any of the completeness criteria.

► **Lemma 5.** Let $\overline{R_i} \in \overline{RDD_{\sigma_i}}, \overline{R'_j} \in \overline{RDD_{\sigma'_j}}$, and denote the representative elements $\overline{\varphi_i}, \overline{\varphi'_j}$. We assume the terms are based only on `map`, `filter` and cartesian without self-products. Let there be fold UDFs $\overline{f_i} : \xi_i \times \sigma_i \rightarrow \xi_i, \overline{f'_j} : \xi'_j \times \sigma'_j \rightarrow \xi'_j$, and initial values $\overline{init_i} : \xi_i, \overline{init'_j} : \xi'_j$. Let there be 2 functions $g : \overline{\xi_i} \rightarrow \xi, g' : \overline{\xi'_j} \rightarrow \xi$. We have: if

$$\bigcup_i FV(\overline{\varphi_i}) \simeq \bigcup_j FV(\overline{\varphi'_j}), \text{ denoted } FV \quad (1)$$

$$g(\overline{init_i}) = g'(\overline{init'_j}) \quad (2)$$

$$\forall \overline{v}, \overline{A_{\varphi_i}}, \overline{A_{\varphi'_j}}. g(\overline{A_{\varphi_i}}) = g'(\overline{A_{\varphi'_j}}) \implies g(f_1(\overline{A_{\varphi_i}}, \varphi_i[\overline{v}/FV])) = g'(f'_1(\overline{A_{\varphi'_j}}, \varphi'_j[\overline{v}/FV])) \quad (3)$$

then $g(\overline{[\varphi_i]_{init_i, f_i}}) = g'(\overline{[\varphi'_j]_{init'_j, f'_j}})$

► **Example 5 (Independent fold).** Below are 2 programs which return a tuple containing the sum of positive elements in its first element, and the sum of negative elements in the second element. We show that by applying lemma 5, we are able to show the equivalence.

Let: $h : (\lambda(P, N), x.ite(x \geq 0, (P + x, N), (P, N - x)))$

	P1($R: RDD_{\text{Int}}$):	P2($R: RDD_{\text{Int}}$):
1	return fold((0, 0), h)(R)	$R_P = \text{filter}(\lambda x. x \geq 0)(R)$
2		$R_N = \text{map}(\lambda x. -x)(\text{filter}(\lambda x. x < 0)(R))$
3		$p = \text{fold}(0, \lambda A, x. A + x)(R_P)$
4		$n = -\text{fold}(0, \lambda A, x. A + x)(R_N)$
5		return (p, n)

$$\Phi(P1) = [\mathbf{x}_R]_{(0,0),h}; \quad \Phi(P2) = ([\phi_{P2}(R_P)]_{0,+}, -[\phi_{P2}(R_N)]_{0,+})$$

$$\phi_{P2}(R_P) = ite(\mathbf{x}_R \geq 0, \mathbf{x}_R, \perp); \quad \phi_{P2}(R_N) = ite(\mathbf{x}_R < 0, -\mathbf{x}_R, \perp)$$

We let $g = \lambda(x, y).(x, y)$ and $g' = \lambda(x, y).(x, -y)$. Induction base case is trivial. Induction step:

$$\forall x, A, B, C. p_1(A) = B \wedge p_2(A) = C \implies h(A, x) = (B + ite(x \geq 0, x, 0), C + ite(x < 0, -x, 0))$$

Substituting for A with B, C , we get that we need to prove:

$$ite(x \geq 0, (B + x, C), (B, C - x)) \stackrel{?}{=} (B + ite(x \geq 0, x, 0), C + ite(x < 0, -x, 0))$$

which is a formula in the Augmented Presburger Arithmetic, whose validity is decidable. ■

5.2.3.2 Nested aggregations.

In the following subsection we present more complex SparkLite programs, in which the value of an aggregate operation is used in later aggregations (i.e. ‘nested’ aggregations). We see that the inductive method is sound in handling those cases, and that under certain conditions, it is complete too.

► **Example 6 (Conditional summation).** The following example takes the *sum* of all elements which are greater than the *count* of elements in an RDD.

Let: $f : (\lambda A, (a, b). A + b)$
 $+: \lambda A, x. A + x$

	P1($R: RDD_{\text{Int}}$):	P2($R: RDD_{\text{Int}}$):
1	$R' = \text{map}(\lambda x.(x, 2))(R)$	$R' = \text{map}(\lambda x.(x, 1))(R)$
2	$sz = \text{fold}(0, f)(R')$	$sz = \text{fold}(0, f)(R')$
3	$B = \text{filter}(\lambda x. x > sz)(R)$	$B = \text{filter}(\lambda x. x > 2 * sz)(R)$
3	return fold(0, +)(B)	return fold(0, +)(B)

The equivalence condition is:

$$\left[\begin{array}{cc} \mathbf{x}_R & \mathbf{x}_R > \phi_{P1}(sz) \\ \perp & \text{otherwise} \end{array} \right]_{0,+} = \left[\begin{array}{cc} \mathbf{x}_R & \mathbf{x}_R > 2 * \phi_{P2}(sz) \\ \perp & \text{otherwise} \end{array} \right]_{0,+}$$

Replacing sz, sz' we get:

$$\left[\begin{array}{cc} \mathbf{x}_R & \mathbf{x}_R > [(\mathbf{x}_R^{(1)}, 2)]_{0,f} \\ \perp & \text{otherwise} \end{array} \right]_{0,+} = \left[\begin{array}{cc} \mathbf{x}_R & \mathbf{x}_R > 2 * [(\mathbf{x}_R^{(1)}, 1)]_{0,f} \\ \perp & \text{otherwise} \end{array} \right]_{0,+}$$

After formally applying lemma 2 we get that the above is equivalent if and only if $\phi_{P1}(sz) = 2 * \phi_{P2}(sz)$, that is:

$$[(\mathbf{x}_R^{(1)}, 2)]_{0,f} = 2 * [(\mathbf{x}_R^{(1)}, 1)]_{0,f}$$

In this case, we set $g = \lambda x.x, g' = \lambda x.2 * x$, and get:

$$0 = g(0) = 2 * 0 \tag{1}$$

$$\begin{aligned} \forall x, A, A'. A = 2 * A' \implies f(A, (x, 2)) &= A + 2 \\ &= 2 * A' + 2 \\ &= 2 * (A' + 1) \\ &= f(A', (x, 1)) \end{aligned} \tag{2}$$

Proving the equivalence. ■

This property is reflected in the following lemma:

► **Lemma 6.** *Let there be two SparkLite programs P_1, P_2 containing two aggregated expressions $a_i = [\varphi_i]_{init_i, f_i}$ for $i \in \{1, 2\}$. Let P_1, P_2 contain return aggregated expressions, the terms of which depend on the previous aggregations a_i : $h_j([\psi_j(a_j)]_{init'_j, g_j})$. If:*

$$FV(\psi_1) \simeq FV(\psi_2), \text{ denoted } FV_\psi \tag{1}$$

$$h_1(init_1) = h_2(init_2) \tag{2}$$

$$\forall \bar{v}, A_1, A_2. h_1(A_1) = h_2(A_2) \implies \tag{3}$$

$$h_1(g_1(A_1, \psi_1(a_1)[\bar{v}/FV_\psi])) = h_2(g_2(A_2, \psi_2(a_2)[\bar{v}/FV_\psi]))$$

Note that Lemma 6 is subject to all previously defined completeness criteria, in conjunction with the completeness criteria applied to the equivalence formula in the induction step. This allows us to define an algorithm for verifying complex equivalences with aggregated queries. The idea is to apply nested inductive proofs (on the nested expressions) during the proof of the induction step of the outer aggregations.

Several examples are given in appendix ??.

References

- 1 Java. <http://java.net>. Accessed: 2016-07-19.
- 2 Python. <https://www.python.org/>. Accessed: 2016-07-19.
- 3 Scala. <http://www.scala-lang.org>. Accessed: 2016-07-19.
- 4 Serge Abiteboul, Richard Hull, and Victor Vianu. *Foundations of Databases*. Addison-Wesley, 1995. URL: <http://webdam.inria.fr/Alice/>.
- 5 Michael Armbrust, Reynold S. Xin, Cheng Lian, Yin Huai, Davies Liu, Joseph K. Bradley, Xiangrui Meng, Tomer Kaftan, Michael J. Franklin, Ali Ghodsi, and Matei Zaharia. Spark sql: Relational data processing in spark. In *Proceedings of the 2015 ACM SIGMOD International Conference on Management of Data*, SIGMOD '15, pages 1383–1394, New York, NY, USA, 2015. ACM. URL: <http://doi.acm.org/10.1145/2723372.2742797>, doi:10.1145/2723372.2742797.
- 6 Aaron R. Bradley and Zohar Manna. *The Calculus of Computation: Decision Procedures with Applications to Verification*. Springer-Verlag New York, Inc., Secaucus, NJ, USA, 2007.
- 7 E. F. Codd. A relational model of data for large shared data banks. *Commun. ACM*, 13(6):377–387, 1970. URL: <http://doi.acm.org/10.1145/362384.362685>, doi:10.1145/362384.362685.
- 8 David C Cooper. Theorem proving in arithmetic without multiplication. *Machine Intelligence*, 1972.
- 9 Ron Cytron, Jeanne Ferrante, Barry K. Rosen, Mark N. Wegman, and F. Kenneth Zadeck. Efficiently computing static single assignment form and the control dependence graph. *ACM Trans. Program. Lang. Syst.*, 13(4):451–490, October 1991. URL: <http://doi.acm.org/10.1145/115372.115320>, doi:10.1145/115372.115320.
- 10 Jeffrey Dean and Sanjay Ghemawat. Mapreduce: Simplified data processing on large clusters. In *Proceedings of the 6th Conference on Symposium on Operating Systems Design & Implementation - Volume 6*, OSDI'04, pages 10–10, Berkeley, CA, USA, 2004. USENIX Association. URL: <http://dl.acm.org/citation.cfm?id=1251254>, doi:10.1145/1251254.1251254.
- 11 Michael J. Fischer and Michael O. Rabin. Super-exponential complexity of presburger arithmetic. Technical report, Massachusetts Institute of Technology, Cambridge, MA, USA, 1974.
- 12 Ross Ihaka and Robert Gentleman. R: A language for data analysis and graphics. *Journal of Computational and Graphical Statistics*, 5(3):pp. 299–314, 1996.
- 13 Viktor Kuncak, Huu Hai Nguyen, and Martin C. Rinard. Deciding boolean algebra with presburger arithmetic. *J. Autom. Reasoning*, 36(3):213–239, 2006.
- 14 Derek C. Oppen. A 222pn upper bound on the complexity of presburger arithmetic. *Journal of Computer and System Sciences*, 16(3):323 – 332, 1978. URL: <http://www.sciencedirect.com/science/article/pii/0022000078900211>, doi:http://dx.doi.org/10.1016/0022-0000(78)90021-1.
- 15 M. Presburger. Über die vollständigkeit eines gewissen systems der arithmetik ganzer zahlen, in welchem die addition als einzige operation hervortritt. *Comptes Rendus du I congrès de Mathématiciens des Pays Slaves*, pages 92–101, 1929.
- 16 Matei Zaharia, Mosharaf Chowdhury, Tathagata Das, Ankur Dave, Justin Ma, Murphy McCauly, Michael J. Franklin, Scott Shenker, and Ion Stoica. Resilient distributed datasets: A fault-tolerant abstraction for in-memory cluster computing. In *Presented as part of the 9th USENIX Symposium on Networked Systems Design and Implementation (NSDI 12)*, pages 15–28, San Jose, CA, 2012. USENIX.
- 17 Matei Zaharia, Mosharaf Chowdhury, Michael J. Franklin, Scott Shenker, and Ion Stoica. Spark: Cluster computing with working sets. In *Proceedings of the 2Nd USENIX Conference on Hot Topics in Cloud Computing*, HotCloud'10, pages 10–10, Berkeley, CA, 2010. USENIX.

USA, 2010. USENIX Association. URL: <http://dl.acm.org/citation.cfm?id=1863103.1863113>.



A Extending Cooper's Algorithm to the Augmented Presburger Arithmetic

► **Proposition 4.** *The theory of formulas over \mathbb{Z}^n with terms in the Augmented Presburger Arithmetic is decidable.*

Proof. Let φ be a quantified formula over $\bigcup_n \mathbb{Z}^n$ with terms in the Augmented Presburger Arithmetic. We shall translate φ to a formula in the Presburger Arithmetic. For any atom $A: a = b$, and $a, b \in \mathbb{Z}^k$ for some $k > 0$, we build the following formula: $\bigwedge_{i=1}^k p_i(a) = p_i(b)$ and replace it in place of A . In the resulting formula, we assign new variable names, replacing the projected tuple variables: For $a \in \mathbb{Z}^k$ we define $x_{a,i} = p_i(a)$ for $i \in \{1, \dots, k\}$. Variable quantification extends naturally, i.e. $\forall a$ becomes $\forall x_{a,1}, \dots, x_{a,k}$, and similarly for \exists . ◀

B Typing rules for SparkLite

Booleans	$\frac{}{\rho \vdash \text{true} : \text{Boolean}}$	$\frac{}{\rho \vdash \text{false} : \text{Boolean}}$
Integers	$\frac{}{\rho \vdash 0, 1, \dots : \text{Integer}}$	
Integer ops	$\frac{\rho \vdash i : \text{Integer}, j : \text{Integer}, \text{op} \in \{+, -, *, \%\}}{\rho \vdash i \text{ op } j : \text{Integer}}$	$\frac{\rho \vdash i : \text{Integer}, j : \text{Integer}, \text{op} \in \{<, \leq, =, \geq, >\}}{\rho \vdash i \text{ op } j : \text{Boolean}}$
Boolean ops	$\frac{\rho \vdash b : \text{Boolean}}{\rho \vdash !b : \text{Boolean}}$	$\frac{\rho \vdash b_1 : \text{Boolean}, b_2 : \text{Boolean}, \text{op} \in \{\wedge, \vee\}}{\rho \vdash b_1 \text{ op } b_2 : \text{Boolean}}$
Tuples	$\frac{\rho \vdash e_1 : \tau_1, e_2 : \tau_2}{\rho \vdash (e_1, e_2) : \tau_1 \times \tau_2}$	$\frac{\rho \vdash e : \tau_1 \times \dots \times \tau_n}{\rho \vdash p_i(e) : \tau_i}$
UDFs	$\frac{\rho \vdash f : C_1 \times \dots \times C_n \rightarrow (\tau \rightarrow \tau'), \bar{e} : C_1 \times \dots \times C_n}{\rho \vdash f(\bar{e}) : \tau \rightarrow \tau'}$	$\frac{\rho \vdash f : \tau \rightarrow \tau', t : \tau}{\rho \vdash f(t) : \tau'}$
RDD	$\frac{\rho \vdash r : \text{RDD}_{\tau}, f : \tau \rightarrow \tau'}{\rho \vdash \text{map}(f)(r) : \text{RDD}_{\tau'}}$	$\frac{\rho \vdash r : \text{RDD}_{\tau}, f : \tau \rightarrow \text{Boolean}}{\rho \vdash \text{filter}(f)(r) : \text{RDD}_{\tau}}$
	$\frac{\rho \vdash r : \text{RDD}_{\tau}, r' : \text{RDD}_{\tau'}}{\rho \vdash \text{cartesian}(r, r') : \text{RDD}_{\tau \times \tau'}}$	$\frac{\rho \vdash r : \text{RDD}_{\tau}, f : \tau' \times \tau \rightarrow \tau, \text{init} : \tau'}{\rho \vdash \text{fold}(\text{init}, f)(r) : \tau'}$

■ **Figure 8** Typing rules for SparkLite

C Proof of Lemma 2

Proof. First we recall the semantics of the **fold** operation on some RDD R , which is a bag. We choose an arbitrary element $a \in R$ and apply the fold function recursively on a and on R with a single instance of a removed. We then write a sequence of elements in the order they are chosen by **fold**: $\langle a_1, \dots, a_n \rangle$, where n is the sum of all multiplicities in the bag R . We also know that a requirement of aggregating operations' UDFs is that they are *commutative*, so the order of elements chosen does not change the final result. We also extend f_i to $\xi_i \times (\sigma_i \cup \{\perp\})$ by setting $f_i(A, \perp) = A$ (\perp is defined to behave as the neutral element for f_i). To prove $g([\varphi_0]_{\text{init}_0, f_0}) = g'([\varphi_1]_{\text{init}_1, f_1})$, it is necessary to prove that

$$g(\llbracket \text{fold} \rrbracket(f_0, \text{init}_0)(R_0)) = g'(\llbracket \text{fold} \rrbracket(f_1, \text{init}_1)(R_1))$$

We set $A_{\varphi_j, 0} = \text{init}_j$ for $j \in \{0, 1\}$. Each element of R_0, R_1 is expressible by providing a concrete valuation to the free variables of φ_0, φ_1 , namely the vector \bar{v} . We choose an arbitrary

sequence of valuations to \bar{v} , denoted $\langle \bar{a}_1, \dots, \bar{a}_n \rangle$, and plug them into the *fold* operation for both R_0, R_1 . The result is 2 sequences of *intermediate values* $\langle A_{\varphi_0,1}, \dots, A_{\varphi_0,n} \rangle$ and $\langle A_{\varphi_1,1}, \dots, A_{\varphi_1,n} \rangle$. We have that $A_{\varphi_j,i} = f_j(A_{\varphi_j,i-1}, \varphi_j[\bar{a}_i/FV])$ for $j \in \{0,1\}$, from the semantics of **fold**. Our goal is to show $g(A_{\varphi_0,n}) = g'(A_{\varphi_1,n})$ for all n . We prove the equality by induction on the *size* of the sequence of possible valuations of \bar{v} , denoted n . In each step i , we show $g(A_{\varphi_0,i}) = g'(A_{\varphi_1,i})$.

Case $n = 0$: $R_0 = R_1 = \emptyset$, so $\llbracket \text{fold} \rrbracket(f_0, \text{init}_0)(R_0) = \text{init}_0$ and $\llbracket \text{fold} \rrbracket(f_1, \text{init}_1)(R_1) = \text{init}_1$. From Equation (3), $g(\text{init}_0) = g'(\text{init}_1)$, as required.

Case $n = i$, assuming correct for $n \leq i-1$: By assumption, we know that the sequence of intermediate values up to $i-1$ is equal up to application of g, g' , and specifically $g(A_{\varphi_0,i-1}) = g'(A_{\varphi_1,i-1})$. We are given the i 'th concrete valuation of \bar{v} , denoted \bar{a}_i . We need to show $A_{\varphi_0,i} = A_{\varphi_1,i}$, so we use the formula for calculating the next intermediate value:

$$\begin{aligned} A_{\varphi_0,i} &= f_0(A_{\varphi_0,i-1}, \varphi_0[\bar{a}_i/FV]) \\ A_{\varphi_1,i} &= f_1(A_{\varphi_1,i-1}, \varphi_1[\bar{a}_i/FV]) \end{aligned}$$

We use Equation (4), plugging in $\bar{v} = \bar{a}_i$, $A_{\varphi_0} = A_{\varphi_0,i-1}$, and $A_{\varphi_1} = A_{\varphi_1,i-1}$. By the induction assumption, $g(A_{\varphi_0,i-1}) = g'(A_{\varphi_1,i-1})$, therefore $g(A_{\varphi_0}) = g'(A_{\varphi_1})$, so Equation (4) yields $g(f_0(A_{\varphi_0}, \varphi_0[\bar{a}_i/FV])) = g'(f_1(A_{\varphi_1}, \varphi_1[\bar{a}_i/FV]))$. By substituting back A_{φ_j} and the formula for the next intermediate value, we get: $g(A_{\varphi_0,i}) = g'(A_{\varphi_1,i})$ as required. \blacktriangleleft