

Verifying Equivalence of Spark Programs

Shelly Grossman¹, Sara Cohen², Shachar Itzhaky³, Noam Rinetzky¹, and Mooly Sagiv¹

- 1 School of Computer Science, Tel Aviv University, Tel Aviv, Israel
`{shellygr,maon,msagiv}@tau.ac.il`
- 2 School of Engineering and Computer Science, The Hebrew University of Jerusalem, Jerusalem, Israel
`sara@cs.huji.ac.il`
- 3 Computer Science and Artificial Intelligence Laboratory, Massachusetts Institute of Technology, USA
`shachari@mit.edu`

Abstract

In this paper, we present a novel approach for verifying the equivalence of Spark programs. Spark is a popular framework for writing large scale data processing applications. Such frameworks, intended for data-intensive operations, share many similarities with traditional database systems, but do not enjoy a similar support of optimization tools. Our goal is to enable such optimizations by first providing the necessary theoretical setting for verifying the equivalence of Spark programs. This is challenging because such programs combine relational algebraic operations with *User Defined Functions (UDFs)*. In this paper, we define a model of Spark as a programming language which imitates Relational Algebra queries in the bag semantics and allows for user defined functions expressible in Presburger Arithmetics. We present a sound and complete reasoning technique for verifying the equivalence of an interesting class of Spark programs.

1 Introduction

The rise of cloud computing and Big Data in the last decade allowed the advent of new programming models, with the intention of simplifying the development process for large-scale needs, letting the programmer focus on business-logic and separating it from the technical details of data management over computer clusters, distribution, communication and parallelization. The first model was MapReduce [13], It allowed programmers to define their logic by composing several iterations of map and reduce operations, where the programmer provided the required map and reduce functions in each step, and the framework was responsible for facilitating the dataflow to the provided functions. MapReduce gave a powerful, yet a clean and abstract programming model. Later on, other frameworks were developed, allowing programmers to write procedural code while still retaining the ability to run it on a large computer cluster, without having the programmer to handle distribution and error recovery. One such framework is Apache Spark [23], in which programmers keep writing code in their programming language of choice, (e.g., Scala [3], Java [1], Python [2], or R [17]) but utilize a special object provided by Spark, called *resilient distributed dataset (RDD)*, providing access to the distributed data itself and to perform transformations on it, using the cloud resources for actual computing. The architecture of Spark comprises of a single master node, referred to as the *driver*, and *worker* nodes in a clustered computer environment. All the nodes have access to the program code, but the driver orchestrates its execution using the underlying Spark framework, abstracting away communications, error recovery, distribution, data partitioning, and parallelization. The access to the data is via the *RDD API*. An *RDD* can be thought of as a simple database table, but which provides support for *User Defined Functions (UDF)*, greatly increasing its expressive power. Spark programs



© Shelly Grossman, Sara Cohen, Shachar Itzhaky, Noam Rinetzky, and Mooly Sagiv;
licensed under Creative Commons License CC-BY

Leibniz International Proceedings in Informatics

LIPIC Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl Publishing, Germany

XX:2 Verifying Equivalence of Spark Programs

handle a family of common tasks involving large datasets, such as log parsing, database queries (via *SparkSQL* [5]), training algorithms and different numeric computations in various fields. Due to this, many Spark programs share several properties: they are mostly short and relatively simple to read and understand. We believe that thanks to this nature of Spark programs, the problem of verifying a program's properties, or even program equivalence as we focus on in this paper, may become feasible, even decidable in an interesting class of Spark programs.

Main Results. The main contributions of this paper can be summarized as follows:

1. We present a simplified model of Spark using a simple programming language called SparkLite, in which UDFs are expressed as simply typed λ -calculus restricted to Presburger Arithmetics. The operations on *RDDs* correspond to operations in relational algebra, with additional aggregate operations expressions, augmented with the added expressivity provided by UDFs.
2. We describe the problem of *program equivalence* (*PE*), and describe a natural class of SparkLite programs, for each *PE* is decidable.
3. We present an algorithm for verifying equivalence of SparkLite programs which is complete for programs in this class.

Find name for class

Related Work. Program and Query Equivalence have long been active research interests. Our model of Spark and its connection to Relational Algebra relies on classical results on Query Equivalence by [?, 22]. In particular, these results were extended to bag semantics, for example in [7, 9, 10, 15]. Our method of verifying aggregate terms using induction has roots in the analysis of programs inductive invariants, for example in . A similar work on combining logic with aggregation can be seen in [16]. For Spark and other similar frameworks, optimization efforts were focused on system optimization and efficient distribution and parallelization, . To our knowledge, there is no general scheme for writing optimizers for data-intensive frameworks such as Spark.

finish

The PL angle?

complete

finish

Overview. Section 2 provides necessary preliminaries for the rest of the paper. In section 3 we give a complete formalization (syntax and semantics) of SparkLite. In section 4 we describe the term semantics of SparkLite problem. In section 5 we present the framework for equivalence checking. In 5.1 we provide a decision procedure for verifying equivalence of SparkLite programs without aggregate expressions. We continue in Section 5.2, where we discuss SparkLite programs with aggregate expressions: we present a sound equivalence verification technique as well as descriptions of a class of SparkLite programs for which the technique can be made complete. In 5.3, we show that the general problem of program equivalence in SparkLite is undecidable. Throughout the paper, missing proofs can be found in the appendix.

2 Preliminaries

In this section, we describe a simple extension of Presburger arithmetic [21], which is the first-order theory of the natural numbers with addition, to tuples of integers, and state its decidability.

Arithmetic Expression	$ae ::= c \mid ae + ae \mid -ae \mid c * ae \mid ae / c \mid ae \% c$
Boolean Expression	$be ::= \text{true} \mid \text{false} \mid e = e \mid ae < ae \mid \neg be \mid be \wedge be \mid be \vee be$
Primitive Expression	$pe ::= ae \mid be$
General Expression	$e ::= pe \mid v \mid (pe, \overline{pe}) \mid p_i(e) \mid \text{ite}(be, e, e)$

■ **Figure 1** Terms of the Augmented Presburger Arithmetic (APA)

Notations. We denote the set of natural numbers (including zero), positive numbers, and integers by \mathbb{N} , \mathbb{N}^+ , and \mathbb{Z} , respectively. We denote a sequence of elements coming from a set X by \overline{X} . We write $\text{ite}(p, e, e')$ to denote an expression which evaluates to e if p holds and to e' otherwise. We use \perp to denote the *undefined* value. A *bag* m over a domain X is a multiset (i.e., a set which allows for repetitions) with elements taken from X . We write $\{\cdot\}$ and $\{\{\cdot\}\}$ to denote sets and bags, respectively. We sometimes use $\{\{;\cdot\}\}$ to denote a bag with explicit multiplicity of the elements. We denote the *size* (number of elements) of a set, respectively, a bag, X by $|X|$.

Presburger Arithmetic. We consider a fragment of first-order logic (FOL) with equality over the integers, where expressions are written in the syntax specified in Figure 1.¹ Disregarding the tuple expressions $((pe, \overline{pe})$ and $p_i(e)$), the resulting first-order theory with the usual \forall and \exists quantifiers is called the *Presburger Arithmetic*. The problem of checking whether a sentence in Presburger arithmetic is valid has long been known to be decidable [14, 21], even when combined with Boolean logic [6, 18], and even with infinities [19].² For example, *Cooper's Algorithm* [11] is a standard decision procedure for Presburger Arithmetic.³⁴

In this paper, we consider a simple extension this language by adding a *tuple constructor* (pe, \overline{pe}) , which allows to create k -tuples, for some $1 \leq k$, of primitive expressions, and a projection operator $p_i(e)$ (projection on the i -th component), which returns the i -th component of a given tuple expression e . and call the extended language *Augmented Presburger Arithmetic* (APA). The decidability of Presburger Arithmetic, as well as Cooper's Algorithm, can be naturally extended to APA. Intuitively, verifying the equivalence of tuple expressions can be done by verifying the equivalence their corresponding constituents.

► **Proposition 1.** *The theory of formulas over \mathbb{Z}^n with terms in APA is decidable.*

3 The SparkLite language

In this section, we define the syntax of SparkLite, a simple imperative programming language which allows to use Spark's *resilient distributed datasets* (RDDs) [23].

3.1 Data Model

Types. SparkLite supports two primitive types: integers (**Int**) and booleans (**Boolean**). On top of this, the user can define types which are Cartesian products of primitive types.

¹ We assume the reader is familiar with FOL, and omit a more formal description for brevity.

² Originally, Presburger Arithmetic was defined as a theory over natural numbers. However, its extension to integers is also decidable. (See, e.g., [6].)

³ The complexity of Cooper's algorithm is $O(2^{2^{2^{pn}}})$ for some $p > 0$ and where n is the number of symbols in the formula [20]. However, in practice, our experiments show that Cooper's algorithm on non-trivial formulas returns almost instantly, even on commodity hardware.

⁴ Basing on [19], we denote infinities $+\infty, -\infty \in \mathbb{Z}$.

In the following we use c to range over integer numerals (constants), $b \in \{\text{true}, \text{false}\}$ to range over Boolean constants, and τ to range over basic types and record types. In addition, SparkLite allows the user to define *RDDs*. RDDs are bags of elements, all of the same type. Hence, RDD_τ denotes bags containing elements of type τ .

Semantic Domains. We interpret the integer and Boolean primitive types as *integers* (\mathbb{Z}) and *booleans* (\mathbb{B}), respectively. We use $\llbracket \cdot \rrbracket$ (semantic brackets) throughout the paper to denote the semantics of program constructs; for types, the semantics is a set of all the values pertaining to it. So $\llbracket \text{Int} \rrbracket = \mathbb{Z}$, $\llbracket \text{Boolean} \rrbracket = \mathbb{B}$.

The interpretation of both primitive types is denoted $T = \mathbb{Z} \cup \mathbb{B}$. The interpretation of all possible types (including mixed Cartesian products of the primitive types) is denoted by $\mathcal{T} = \bigcup_n T^n$.

An *RDD* type is interpreted as a *bag* (an unordered set allowing repeating elements). We write $\llbracket RDD_\tau \rrbracket = (\llbracket \tau \rrbracket \leftrightarrow \mathbb{N})$, meaning that an RDD value r of type RDD_τ is interpreted as a bag of values from τ , that is $\llbracket r \rrbracket \in (\llbracket \tau \rrbracket \leftrightarrow \mathbb{N})$. We let $\llbracket RDD \rrbracket = \bigcup_{\tau \in \mathcal{T}} \llbracket RDD_\tau \rrbracket$, the semantic domain of RDDs over all possible record types $\tau \in \mathcal{T}$.

3.2 Functional Model

Operations. *RDDs* are analogous to database tables and as such the methods to query the *RDDs* are inspired by both *Relational Algebra* (*RA*) [4, 8] and Spark [24]. *RA* has 5 basic operators, which are *Select*, *Project*, *Cartesian Product*, *Union* and *Subtract*. This paper focuses on the first three operators. The *Select* operator is analogous to *filter* in SparkLite, and *Project* is analogous to *map*. The expressive power of SparkLite’s *map* and *filter* is greater than their analogous *RA* operations thanks to UDFs (see next), which allow *extended projection* as well as greater flexibility in executing complex operations on elements of different types.

UDFs. A special feature of Spark is allowing some of its standard operations to be *higher order functions* — they take a function and apply it to an RDD in a specific manner defined by the operation. For example, the operation *fold* can be applied to an RDD containing integer numbers by providing a function that adds two integers, yielding the sum of all the numbers in the bag. Such a function is called a “*User-Defined Function*”, or *UDF*, for short. The *signature* of a UDF contains information on the return type and the arguments types, and when applied in the context of an RDD operation, the signature should match both the RDD type and the operation on it (see typing rules in appendix B). Each UDF has a *definition*, which takes the syntax $\lambda \bar{v}. e$, where \bar{v} are names of variables used as arguments. For example, $\text{addMod10} = \lambda x, y. x \% 10 + y \% 10$. This function’s *signature* would be $\text{addMod10}: \text{Int} \times \text{Int} \rightarrow \text{Int}$, that is, it takes two *Int* arguments and produces one *Int* value. The types are usually omitted from the definition for brevity, but when the types are not clear from the context, we may write them as annotations: $\text{addMod10} = \lambda x: \text{Int}, y: \text{Int}. x \% 10 + y \% 10: \text{Int}$. We could also write it as a function with single argument like this: $\text{addPairMod10} = \lambda z. p_1(z) \% 10 + p_2(z) \% 10$, where it would have the signature $\text{addPairMod10}: (\text{Int} \times \text{Int}) \rightarrow \text{Int}$. For readability, we sometimes give names to the projections and write them implicitly as $\text{addPairMod10} = \lambda(x, y). x \% 10 + y \% 10$; this is just a shortcut.

We allow the definition of these functions to be *parametric*, by using two levels of λ . The outer level denotes the parameters, which can be provided to obtain a regular function. For example, a function that adds 1 to an integer is written as $f = \lambda x. x + 1$, whereas a

Basic Types	τ	$::=$	$\text{int} \mid \text{bool} \mid \tau \times \dots \times \tau$
RDDs	RDD	$::=$	RDD_τ
Variables	x	$::=$	$v \mid r$
Arithmetic Exp.	ae	$::=$	$c \mid ae + ae \mid -ae \mid c * ae \mid ae / c \mid ae \% c$
Boolean Exp.	be	$::=$	$\text{true} \mid \text{false} \mid e = e \mid ae < ae \mid \neg be \mid be \wedge be \mid be \vee be$
General Basic Exp.	e	$::=$	$ae \mid be \mid v \mid (e, e) \mid p_i(e) \mid \text{if } (b) \text{ then } e \text{ else } e$
Functions	$Fdef$	$::=$	$\text{def } f = \lambda \overline{y}:\overline{\tau} \ e:\tau$
Parametric Functions	$PFdef$	$::=$	$\text{def } F = \lambda \overline{x}:\overline{\tau}. \lambda \overline{y}:\overline{\tau} \ e:\tau$
RDD Exp.	re	$::=$	$\text{cartesian}(r, r) \mid \text{map}(f)(r) \mid \text{filter}(f)(r)$
RDD Aggregation Exp.	ge	$::=$	$\text{fold}(e, f)(r)$
General Exp.	η	$::=$	$e \mid re \mid ge$
Program Body	E	$::=$	$\text{Let } x = \eta \text{ in } E \mid \eta$
Program	$Prog$	$::=$	$P(\overline{r}:RDD_\tau, \overline{v}:\overline{\tau}) = \overline{Fdef} \ \overline{PFdef} \ E$

■ **Figure 2** Syntax for SparkLite

function that adds any constant is written as $g = \lambda a. \lambda x. x + a$. The function is *curried*, so that applying it to parameter values of the appropriate types produces a function (by *beta-reduction*): $g(1)$ is identical to f .

3.3 Syntax

The syntax of SparkLite language is defined in Figure 2.

Syntactic Categories. We assume variables to be an infinite syntactic category, ranged over by $v, b, r \in \text{Vars}$. Expressions range over e . An integer constant is denoted c . There are 4 operations: `map`, `filter`, `cartesian`, and `fold`. Some of the operations require arguments, which may be either a primitive expression, an RDD , or a function. Functions range over $f, F \in \text{LambdaExpressions}$. Parametric functions are denoted by capital meta-variables (F as opposed to f for regular functions) and must always be given the list of parameters when passed to an operation.

Program structure. The header of a program contains function definitions. Loops are not allowed in the body of a program. Variable declarations are in *SSA* (*Static Single Assignment*) form [12]. Variables are immutable by this construction. Programs have no side effects, do not change the inputs, and always return a value. The *program signature* will consist of its name, its input types and return type: $P(\overline{\tau}_i, \overline{RDD}_i):\mathcal{T}_o$

Example program. Consider the example SparkLite program in Figure 3. From the example program we can see the general structure of SparkLite programs: First, the functions that are used as UDFs in the program are declared and defined: `isOdd`, `sumFlatPair` defined as $Fdef$, and `doubleAndAdd` defined as a $PFdef$. The name of the program ($P = P1$) is announced with a list of input RDDs (R_0, R_1) ($Prog$ rule). Instead of writing *Let* l_1 *in* *Let* l_2 *in* ..., we use syntactic sugar, where each line of code contains a single l_i , and the last line denotes the

$$isOdd = \lambda x: \text{Int}. \neg(x \% 2 = 0)$$

$$\text{Let: } doubleAndAdd = \lambda c: \text{Int}. \lambda x: \text{Int}. 2 * x + c$$

$$sumFlatPair = \lambda A: \text{Int}, (x, y): \text{Int} \times \text{Int}. A + x + y$$

	$P1(R_0: RDD_{\text{Int}}, R_1: RDD_{\text{Int}}):$
1	$A = \text{filter}(isOdd)(R_0)$
2	$B = \text{map}(doubleAndAdd(1))(A)$
3	$C = \text{cartesian}(B, R_1)$
4	$v = \text{fold}(0, sumFlatPair)(C)$
5	return v

■ **Figure 3** Example SparkLite program

return value using the **return** keyword. Here, 3 variables of RDD type (A, B, C) and one integer variable (v) are bound by *Lets*. We can see in the definition of A an application of the *filter* operation, accepting the RDD R_0 and the function *isOdd*. For B 's definition we apply the *map* operation with a parametric function *doubleAndAdd* with the parameter 1, which is interpreted as $\lambda x. 2 * x + 1$. C is the cartesian product of B and input RDD R_1 . We apply an aggregation using *fold* on the RDD C , with an initial value 0 and the function *sumFlatPair*, which 'flattens' elements of tuples in C , taking their sum. The sum total of all this elements is stored in the variable v . The returned value is the integer variable v . The program's signature is $P1(RDD_{\text{Int}}, RDD_{\text{Int}}): \text{Int}$.

3.4 Operational Semantics

Program Environment. We define a unified semantic domain $\mathcal{D} = \mathcal{T} \cup RDD$ for all types in SparkLite. The *program environment* type: $\mathcal{E} = \text{Vars} \rightarrow \mathcal{D}$ is a mapping from each variable in **Vars** to its value, according to type. A variable's type does not change during the program's run, nor does its value.

Data flow. We start with an initial environment function ρ_0 maps all input variables and function definitions. We define the *semantic interpretation* of expressions based on an environment $\rho \in \mathcal{E}$, and specifically for $x \in \bar{\tau} \cup \bar{\nu}$, $\llbracket x \rrbracket(\rho) = \rho(x)$. The semantics of composite expressions are straight-forward using the semantics of their components. The semantics of *Let* is to create a new environment by binding the variable name. In Figure 4 we specify the behavior of $\llbracket \cdot \rrbracket(\cdot)$ for all expressions and statements.

Notes. In *map*, if f maps y to x , the multiplicity of x is the sum of multiplicities of all y elements. In other words, if an element x appears n times, we apply the f on it n times. *fold* is *well defined*: A fold UDF f should satisfy $\forall A, x, y. f(f(A, x), y) = f(f(A, y), x)$ (order of elements does not affect the result).

Example. For the example program Figure 3, suppose we were given the following input: $R_0 = \{(1; 7), (2; 1)\}$, $R_1 = \{(3; 4), (5; 2)\}$. Then: $\rho(A) = \{(1; 7)\}$, $\rho(B) = \{(3; 7)\}$, $\rho(C) = \{((3, 3); 28), ((3, 5); 14)\}$, $\rho(v) = 28 * (3+3) + 14 * (3+5)$ and the program returns $\rho(v) = 280$.

$$\begin{aligned}
\llbracket c \rrbracket(\rho) &= c \\
\llbracket v \rrbracket(\rho) &= \rho(v) \\
\llbracket \text{unOp } e \rrbracket(\rho) &= \text{unOp } \llbracket e \rrbracket(\rho) \\
\llbracket e_1 \text{ binOp } e_2 \rrbracket(\rho) &= \llbracket e_1 \rrbracket(\rho) \text{ binOp } \llbracket e_2 \rrbracket(\rho) \\
\llbracket (e_1, \dots, e_n) \rrbracket(\rho) &= (\llbracket e_1 \rrbracket(\rho), \dots, \llbracket e_n \rrbracket(\rho)) \\
\llbracket \text{if } e_1 \text{ then } e_2 \text{ else } e_3 \rrbracket(\rho) &= \text{ite}(\llbracket e_1 \rrbracket(\rho), \llbracket e_2 \rrbracket(\rho), \llbracket e_3 \rrbracket(\rho)) \\
\llbracket \text{map}(f)(r) \rrbracket(\rho) &= \{\{\rho(f)(x) \mid x \in \rho(r)\}\} \\
\llbracket \text{filter}(b)(r) \rrbracket(\rho) &= \{\{x \mid x \in \rho(r) \wedge \rho(b)(x)\}\} \\
\llbracket \text{cartesian}(r_1, r_2) \rrbracket(\rho) &= \{\{(x_1, x_2) \mid x_1 \in \rho(r_1) \wedge x_2 \in \rho(r_2)\}\} \\
\llbracket \text{fold}(a_0, f)(r) \rrbracket(\rho) &= q(\llbracket a_0 \rrbracket(\rho), \rho(r)), \text{ where} \\
&\quad q(v_0, s) = \begin{cases} v_0 & s = \emptyset \\ \rho(f)(x, q(v_0, s')) & s = \{\{x; 1\}\} \cup s' \end{cases} \\
\llbracket \text{Let } x = \eta \text{ in } E \rrbracket(\rho) &= \llbracket E \rrbracket(\rho[x \mapsto \llbracket \eta \rrbracket(\rho)]) \\
\llbracket P(\dots) = \dots E \rrbracket(\rho_0) &= \llbracket E \rrbracket(\rho_0)
\end{aligned}$$

■ **Figure 4** Semantics of SparkLite. $\text{Prog} = P(\overline{r} : \overline{RDD\tau}, \overline{v} : \overline{\tau}) = \overline{Fdef} \ \overline{PFdef} \ E$. unOp and binOp are taken from Figure 2: $\text{unOp} \in \{-, \neg, \pi_i\}$, $\text{binOp} \in \{+, *, /, \%, =, <, \wedge, \vee, (,)\}$

$$\begin{aligned}
\phi_P(\text{Let } x = \eta \text{ in } E) &= \phi_P(E')[\phi_P(\eta)/x] \\
\phi_P(e) &= e \\
\phi_P(\text{map}(f)(r)) &= f(\phi_P(r)) \\
\phi_P(\text{filter}(f)(r)) &= \text{ite}(f(\phi_P(r)) = \text{tt}, \phi_P(r), \perp) \\
\phi_P(\text{cartesian}(r_1, r_2)) &= (\phi_P(r_1), \phi_P(r_2)) \\
\phi_P(\text{fold}(f, e)(r)) &= [\phi_P(r)]_{e, f} \\
\phi_P(r) &= \begin{cases} \mathbf{x}_r & r \in \overline{r} \\ r & \text{otherwise} \end{cases}
\end{aligned}$$

$$\begin{aligned}
\text{Let } P : P(\overline{r}, \overline{v}) &= \overline{F} \ \overline{f} \ E \\
\Phi(P) &= \phi_P(E)
\end{aligned}$$

■ **Figure 5** Compiling SparkLite to logical terms (ϕ).

$$\begin{aligned}
\llbracket \mathbf{x}_{r_i} \rrbracket(\overline{v}, \overline{r}) &= r_i \\
\llbracket f(t) \rrbracket(\overline{v}, \overline{r}) &= \{\{\llbracket f \rrbracket(z) \mid z \in \llbracket t \rrbracket(\overline{v}, \overline{r})\}\} \\
\llbracket \text{ite}(f(t), t, \perp) \rrbracket(\overline{v}, \overline{r}) &= \{\{z \mid z \in \llbracket t \rrbracket(\overline{v}, \overline{r}) \wedge \llbracket f \rrbracket(z)\}\} \\
\llbracket (t, t') \rrbracket(\overline{v}, \overline{r}) &= \{\{(z, z') \mid z \in \llbracket t \rrbracket(\overline{v}, \overline{r}) \wedge z' \in \llbracket t' \rrbracket(\overline{v}, \overline{r})\}\} \\
\llbracket [t]_{e, f} \rrbracket(\overline{v}, \overline{r}) &= \llbracket \text{fold}(e, f)(r') \rrbracket[r' \mapsto \llbracket t \rrbracket(\overline{v}, \overline{r})]
\end{aligned}$$

■ **Figure 6** Semantics of terms.

4 Term Semantics for SparkLite

In this section, we present an alternative, equivalent semantics for SparkLite where the program is interpreted as a term in APA. This term is called the *program term* and denoted $\Phi(P)$ for program P , specified in Figure 5. These terms are assigned their own semantics such that the semantics of $\Phi(P)$ is identical to the semantics of P as defined earlier. Special variables are used to refer to elements of input RDDs⁵, and a new language construct is added to denote the fold operation. As an example, we take the program $P1$ from Figure 3, and show how to construct $\Phi(P1)$ by recursively applying ϕ_P and simplifying the formula.

$$\begin{aligned}
\phi_{P1}(A) &= \phi_{P1}(\text{filter}(\text{isOdd})(R_0)) = \text{ite}(\text{isOdd}(\phi_{P1}(R_0)), \phi_{P1}(R_0), \perp) \\
&= \text{ite}(\text{isOdd}(\mathbf{x}_{R_0}), \mathbf{x}_{R_0}, \perp) \\
\phi_{P1}(B) &= \phi_{P1}(\text{doubleAndAdd}(1)(A)) = \text{doubleAndAdd}(1)(A) \\
&= \text{doubleAndAdd}(1)(\text{ite}(\text{isOdd}(\mathbf{x}_{R_0}), \mathbf{x}_{R_0}, \perp)) \\
\phi_{P1}(C) &= \phi_{P1}(\text{cartesian}(B, R_1)) = (B, \mathbf{x}_{R_1}) \\
\phi_{P1}(v) &= \phi_{P1}(\text{fold}(0, \text{sumFlatPair})(C))[C]_{0, \text{sumFlatPair}} \\
\Phi(P1) &= \phi_P(v) = [C]_{0, \text{sumFlatPair}} = [(B, \mathbf{x}_{R_1})]_{0, \text{sumFlatPair}} \\
&= [(\text{doubleAndAdd}(1)(\text{ite}(\text{isOdd}(\mathbf{x}_{R_0}), \mathbf{x}_{R_0}, \perp)), \mathbf{x}_{R_1})]_{0, \text{sumFlatPair}}
\end{aligned}$$

Representative elements of RDDs. The variables assigned by ϕ for input RDDs are called *representative elements*. In a program that receives an input RDD r , we denote the representative element of r as: \mathbf{x}_r . The set of possible valuations to that variable is equal to the bag defined by r , and an additional ‘undefined’ value (\perp), for the empty RDD. Therefore \mathbf{x}_r ranges over $\text{dom}(r) \cup \{\perp\}$. By abuse of notation, the term for a non-input RDD, computed in a SparkLite program, is also called a representative element.

Formalization of the Term Semantics for SparkLite. Let P be a SparkLite program. We use standard notations \bar{r} for the inputs of P , and r^{out} for the output of P . The term $\Phi(P)$ is called the *program term of P* as before. The *Term Semantics* (TS) of a program that returns an RDD-type output is the bag that is obtained from all possible valuations to the free variables:

$$TS(P)(\bar{v}, \bar{r}) = \llbracket \Phi(P) \rrbracket(\bar{v}, \bar{r})$$

Where the meaning of $\llbracket \Phi(P) \rrbracket$ is determined according to Figure 6. Assigning a concrete valuation to the free variables of $\Phi(P)$ returns an element in the output RDD r^{out} . By taking all possible valuations to the term with elements from \bar{r} , we get the bag equal to r^{out} .

► **Proposition 2.** *Let $P : P(\bar{r}) = \bar{F} \bar{f} E$ be a SparkLite program, $\llbracket P \rrbracket$ be the interpretation of its output according to the operational semantics, and $TS(P)$ by the term semantics of P . Then, for any input \bar{v}, \bar{r} , we have:*

$$TS(P)(\bar{v}, \bar{r}) = \llbracket P \rrbracket(\llbracket \bar{v} \rrbracket, \llbracket \bar{r} \rrbracket)$$

5 Verifying Equivalence of SparkLite Programs

The Program Equivalence (PE) problem. Let P_1 and P_2 be SparkLite programs, with signature $P_i(\bar{T}, \overline{RDD_T}) : \tau$ for $i \in \{1, 2\}$. We use $\llbracket P_i \rrbracket(\llbracket \bar{v} \rrbracket, \llbracket \bar{r} \rrbracket)$ to denote the result of P_i .

⁵ To avoid overhead of notations, we assume the programs do not contain self-products (for every product in the program, the sets of variables appearing in each component must be disjoint).

We say that P_1 and P_2 are *equivalent*, if for all input values \bar{v} and RDDs \bar{r} , it holds that $\llbracket P_1 \rrbracket(\llbracket \bar{v} \rrbracket, \llbracket \bar{r} \rrbracket) = \llbracket P_2 \rrbracket(\llbracket \bar{v} \rrbracket, \llbracket \bar{r} \rrbracket)$.

5.1 Verifying Equivalence of SparkLite Programs without Aggregations

Given two programs P_1, P_2 receiving as input a series of RDDs $\bar{r} = (r_1, \dots, r_n)$. We assume w.l.o.g. the programs do not receive non-RDD arguments \bar{v} .⁶ We let $\bar{x} = (x_1, \dots, x_n)$ be a concrete valuation for all input RDDs representative elements: \mathbf{x}_{r_i} will map to x_i . We denote the substitution of the concrete valuation in a term t over $\bar{\mathbf{x}}_r$ as: $t(\bar{x}) = t[x_1/\mathbf{x}_{r_1}, \dots, x_n/\mathbf{x}_{r_n}]$.

Comparing representative elements. For two program terms to be comparable, they must depend on the same input RDDs. For example:

	$P1(R_0: RDD_{\text{Int}}, R_1: RDD_{\text{Int}}):$	$P2(R_0: RDD_{\text{Int}}, R_1: RDD_{\text{Int}}):$
1	<code>return map($\lambda x.1$)(R_0)</code>	<code>return map($\lambda x.1$)(R_1)</code>

$P1$ and $P2$ have the same program term (the constant 1), but the multiplicity of that element in the output bag is different and depends on the source input RDD. In $P1$, its multiplicity is the same as the size of R_0 , and in $P2$ it is the same as the size of R_1 . $P1$ and $P2$ are therefore not equivalent, because we can provide inputs R_0, R_1 of different sizes. Therefore, for each program term $\Phi(P)$ we consider the *set of free variables*, $FV(\Phi(P))$. Each free variable has some source input RDD. In the example, $FV(\Phi(P1)) = \{\mathbf{x}_{R_0}\}$, and $FV(\Phi(P2)) = \{\mathbf{x}_{R_1}\}$.

► **Proposition 3.** *Let there be two terms t_1, t_2 , over input RDDs \bar{r} . such that $FV(t_1) \neq FV(t_2)$, and $t_1 \neq \perp \vee t_2 \neq \perp$. Then $\exists \bar{r}. \llbracket t_1 \rrbracket(\bar{r}) \neq \llbracket t_2 \rrbracket(\bar{r})$.*

The program terms may contain \perp expressions, therefore we need to encode the formulas in APA, and remove all appearances of \perp , which is not part of it. We write a series of universally true schemes for translating terms referencing \perp to APA when appearing in an equivalence formula, including translation of all conditionals to FOL. Iterative application of these rules transforms the equivalence formula to a formula in APA, without *ite* and without \perp .

► **Proposition 4** (Schemes for converting conditionals to a normal form). *Let t denote terms without *ite* constructs in them, c are conditions, and f, g are functions which are extended to return \perp if one of the given arguments is \perp .*

1. *The following useful identity for applying functions on a conditional is true:*

$$f(\text{ite}(c, t_1, t_2)) = \text{ite}(c, f(t_1), f(t_2))$$

2. *Equivalence of functions of conditionals:*

$$(f(\text{ite}(c, t, \perp)) = g(\text{ite}(c', t', \perp))) \iff ((c \iff c') \wedge (c \implies f(t) = g(t')))$$

3. *Equivalence of a function of a conditional and an arbitrary term without conditionals:*

$$(f(\text{ite}(c, t, \perp)) = t') \iff (c \wedge f(t) = t')$$

⁶ The extension to equivalence of terms based also on non-RDD inputs is immediate by quantification on the non-RDD variables in the term.

XX:10 Verifying Equivalence of Spark Programs

1. If $\Phi(P_1) = \perp \wedge \Phi(P_2) = \perp$, output **equivalent**.
2. Verify that: $FV(\Phi(P_1)) = FV(\Phi(P_2))$. If not, output **not equivalent**.
3. We check the following formula is satisfiable:

$$\exists \bar{x}. \Phi(P_1)[\bar{v}/FV(\Phi(P_1))] \neq \Phi(P_2)[\bar{v}/S(FV(\Phi(P_2)))]$$

If it is satisfiable, return **not equivalent**. Otherwise, the formula is unsatisfiable, return **equivalent**.

■ **Figure 7** An algorithm for solving *PE* for two programs P_1, P_2 with the same signature

4. *Applying a function with multiple arguments on conditionals and terms without conditionals:*

$$f(ite(c_1, t_1, \perp), \dots, ite(c_n, t_n, \perp), t'_1, \dots, t'_m) = ite(\bigwedge_{i=1}^n c_i, f(t_1, \dots, t_n, t'_1, \dots, t'_m), \perp)$$

5. *Unnesting of nested conditionals:*

$$ite(c_{ext}, ite(c_{int}, t, \perp), \perp) = ite(c_{int} \wedge c_{ext}, t, \perp)$$

6. *General conditionals:*

$$(ite(c, t_1, t_2) = ite(c', t'_1, t'_2)) \iff ((c \wedge c' \implies t_1 = t'_1) \wedge (c \wedge \neg c' \implies t_1 = t'_2) \wedge (\neg c \wedge c' \implies t_2 = t'_1) \wedge (\neg c \wedge \neg c' \implies t_2 = t'_2))$$

7. *Comparison to \perp (reminder: t does not contain ite):*

$$(ite(c, t, \perp) = \perp) \iff \neg c$$

► **Theorem 1** (Decidability for programs without aggregations). *Given two SparkLite programs P_1 and P_2 which do not contain aggregate operations, *PE* is decidable.*

Proof. For non-RDD return types, the absence of aggregate operators implies we can use Proposition 1, as the returned expression is expressible in APA. For RDDs we provide an algorithm in Figure 7, which is a decision procedure: If both program terms result in an empty bag (step 1), the algorithm detects it and outputs the programs are equivalent. Otherwise, the algorithm checks syntactically in step 2 that $FV(\Phi(P_1)) = FV(\Phi(P_2))$, and outputs the programs are not equivalent if that is not the case - as justified by Proposition 3. The correctness of the algorithm in step 3 follows from the equivalence of the TS semantics and the operational semantics defined in 3.4 (Proposition 2). The equivalence formula generated does not contain aggregate terms, and applications of *ite* are normalized using the rules in Proposition 4, resulting in a formula definable in APA. The algorithm generates formulas in APA several times: once to verify whether the programs do not return the empty bag for all inputs, and second to test for equivalence.⁷ From Proposition 1, all the formulas checked by the algorithm are decidable. ◀

⁷ In the next section, program terms may contain aggregate expressions. In that case, there may be more formulas generated, and subsequently more calls to Cooper's Algorithm.

Examples. Note: All examples use syntactic sugar for ‘*Let*’ expressions. For brevity, instead of applying ϕ on the underlying ‘*Let*’ expressions, we apply it line-by-line from the top-down. In addition, we assume that in programs returning an RDD-type, the RDD is named r^{out} , and the programs always end with **return** r^{out} . Thus, $\Phi(P) = \phi_P(r^{out})$.

► **Example 1 (Basic optimization - operator pushback).** This example shows a common optimization of pushing the filter/selection operator backward, to decrease the size of the dataset.

	P1($R: RDD_{Int}$):	P2($R: RDD_{Int}$):
1	$R' = \text{map}(\lambda x. 2 * x)(R)$	$R' = \text{filter}(\lambda x. x < 7)(R)$
2	return $\text{filter}(\lambda x. x < 14)(R')$	return $\text{map}(\lambda x. x + x)(R')$

Both programs may return an non-empty RDD of integers, and the sets of free variables are equal: $FV(\Phi(P1)) = \{\mathbf{x}_R\} = FV(\Phi(P2))$. We analyze the representative elements:

$$\begin{aligned} \phi_{P1}(R') &= 2 * \mathbf{x}_R; & \phi_{P1}(r^{out}) &= \text{ite}(\varphi < 14 \wedge \varphi = \phi_{P1}(R'), \varphi, \perp) = \text{ite}(2 * \mathbf{x}_R < 14, 2 * \mathbf{x}_R, \perp) \\ \phi_{P2}(R') &= \text{ite}(\mathbf{x}_R < 7, \mathbf{x}_R, \perp); & \phi_{P2}(r^{out}) &= (\lambda x. x + x)(\phi_{P1}(R')) = \text{ite}(\mathbf{x}_R < 7, \mathbf{x}_R, \perp) + \text{ite}(\mathbf{x}_R < 7, \mathbf{x}_R, \perp) \end{aligned}$$

We need to verify that:

$$\forall \mathbf{x}_R. \text{ite}(2 * \mathbf{x}_R < 14, 2 * \mathbf{x}_R, \perp) = \text{ite}(\mathbf{x}_R < 7, \mathbf{x}_R, \perp) + \text{ite}(\mathbf{x}_R < 7, \mathbf{x}_R, \perp)$$

Using Proposition 4, $\text{ite}(2 * \mathbf{x}_R < 14, 2 * \mathbf{x}_R, \perp) = \text{ite}(\mathbf{x}_R < 7, \mathbf{x}_R, \perp) + \text{ite}(\mathbf{x}_R < 7, \mathbf{x}_R, \perp)$ becomes via rule (4): $\text{ite}(2 * \mathbf{x}_R < 14, 2 * \mathbf{x}_R, \perp) = \text{ite}(\mathbf{x}_R < 7 \wedge \mathbf{x}_R < 7, \mathbf{x}_R + \mathbf{x}_R, \perp)$ which via rule (2) becomes: $(2 * \mathbf{x}_R < 14 \iff \mathbf{x}_R < 7) \wedge (2 * \mathbf{x}_R < 14 \implies 2 * \mathbf{x}_R = \mathbf{x}_R + \mathbf{x}_R)$. See ?? for an implementation. ■

For additional examples, refer to appendix ??.

5.2 Verifying Equivalence of a Class of SparkLite Programs with Aggregation

In the following section we discuss how the existing framework can be extended to prove equivalence of SparkLite programs containing aggregate expressions. The terms for aggregate operations are given using a special operator - $[t]_{i,f}$ - where t is the term being folded, i is the initial value, and f is the fold function. The operator *binds* all free variables in the term t , thus the free variables of t are not contained in the free variables set of the term that includes the aggregate term.

5.2.1 Single aggregate

The simplest class of programs in which an aggregation operator appears, is programs whose program terms are a function of an aggregate term, that is have the form $g([t]_{i,f})$.

► **Definition 1** (The Agg^1 class). *Let there be a program P with $\Phi(P) = g([t]_{i,f})$. $P \in Agg_R^1$ if t does not contain aggregate terms.*

► **Lemma 2** (Sound method for verifying equivalence of Agg^1 programs). *Let there be representative elements φ_0, φ_1 over σ_0, σ_1 . Let there be two fold functions $f_0 : \xi_0 \times \sigma_0 \rightarrow \xi_0, f_1 :$*

XX:12 Verifying Equivalence of Spark Programs

$\xi_1 \times \sigma_1 \rightarrow \xi_1$, two initial values $init_0: \xi_0, init_1: \xi_1$, and two functions $g: \xi_0 \rightarrow \xi, g': \xi_1 \rightarrow \xi$. We have $g([\varphi_0]_{init_0, f_0}) = g'([\varphi_1]_{init_1, f_1})$ if:

$$FV(\varphi_0) = FV(\varphi_1) \quad (1)$$

$$g(init_0) = g'(init_1) \quad (2)$$

$$\begin{aligned} \forall \bar{v}, A_{\varphi_0}: \xi_0, A_{\varphi_1}: \xi_1. g(A_{\varphi_0}) = g'(A_{\varphi_1}) \implies \\ g(f_0(A_{\varphi_0}, \varphi_0(\bar{v}))) = g'(f_1(A_{\varphi_1}, \varphi_1(\bar{v}))) \end{aligned} \quad (3)$$

Lemma 2 shows that an inductive proof of the equality of folded values is *sound*. Therefore, given two folded expressions which are not equivalent, the lemma is guaranteed to report so.

► **Example 2 (Maximum and minimum).** Below is an example of two equivalent programs belonging to Agg^1 :

Let: $max = \lambda M, x. \text{if}(x > M) \text{ then } \{x\} \text{ else } \{M\}$ $min = \lambda M, x. \text{if}(x < M) \text{ then } \{x\} \text{ else } \{M\}$		
	P1($R: RDD_{Int}$):	P2($R: RDD_{Int}$):
1	return $\text{fold}(-\infty, max)(R)$	$R' = \text{map}(\lambda x. -x)(R)$
2		return $-\text{fold}(+\infty, min)(R')$

The programs compute the maximum element of a numeric RDD in two different methods: in the first program by getting the maximum directly, and in the second by getting the additive inverse of the minimum of the additive inverses of the elements. The equivalence formula is:

$$[\mathbf{x}_R]_{-\infty, max} = -[-\mathbf{x}_R]_{+\infty, min}$$

We apply Lemma 2: The two program apply a fold operation on a term of the same RDD R . $init_0 = -\infty, init_1 = +\infty$ and $g = \lambda x.x., g' = \lambda x. -x$, therefore $g(-\infty) = g'(+\infty)$ as required. We check the inductive claim:

$$\forall x, A, A'. A = -A' \implies max(A, \mathbf{x}_R(x)) = -min(A', -\mathbf{x}_R(x))$$

We assume $A = -A'$ and attempt to prove $max(A, x) = -min(A', -x)$, after normalizing to APA:

$$\begin{aligned} max(A, x) &= ? -min(A', -x) \\ ite(A > x, A, x) &= ? -ite(A' < -x, A', -x) = ite(A' < -x, -A', x) \end{aligned}$$

And we verify the following APA formula:

$$\begin{aligned} \forall x, A, A'. A = -A' \implies ((A > x \wedge A' < -x \implies A = -A') \wedge (A > x \wedge A' \geq -x \implies A = x) \\ \wedge (A \leq x \wedge A' < -x \implies x = -A') \wedge (A \leq x \wedge A' \geq -x \implies x == x)) \end{aligned}$$

which is true. ■

Additional Examples. Refer to appendix ??.

Completeness There are several cases in which one or more of the requirements of Lemma 2 are not satisfied, yet the aggregate expressions are equal. The first requirement, $FV(\varphi_0) = FV(\varphi_1)$, is not necessary when the fold applied on the terms are both *trivial*.

► **Definition 2 (Trivial fold).** $[\varphi]_{init, f}$ is a trivial fold if:

$$\forall \bar{v}. f(init, \varphi(\bar{v})) = init$$

If two instances of Agg^1 have trivial folds, then Equation (2) in Lemma 2 is a sufficient condition for the equivalence:

$$g([\varphi_0]_{init_0, f_0}) = g(init_0) \wedge g'([\varphi_1]_{init_1, f_1}) = g'(init_1) \wedge g(init_0) = g'(init_1) \implies g([\varphi_0]_{init_0, f_0}) = g'([\varphi_1]_{init_1, f_1})$$

Conversely, when the fold is not trivial, the proof of Lemma 2 requires the sets of free variables to be isomorphic. Otherwise, the induction termination is not well defined. One possibility is that the size of participating RDDs may not be equal. Assuming one set of free variables is contained in the other, any non-constant result of the fold function on these additional elements will lead to unequal results. To avoid such peculiarities, we shall require for additional classes of programs to satisfy equal sets of free variables in their aggregate terms. We proceed with an example showing a case which Lemma 2 does not cover.

► **Example 3 (Non-injective modification of folded expressions).** Non-injective transformations can weaken the inductive claim, resulting in failure to prove it. As a result, Lemma 2 fails to prove the equivalence of the following two Agg_1 programs.

	P1($R: RDD_{Int}$):	P2($R: RDD_{Int}$):
1	$R' = \text{map}(\lambda x.x\%3)(R)$	$v = \text{fold}(0, \lambda A, x.A + x)(R)$
2	$\text{return fold}(0, \lambda A, x.(A + x)\%3)(R') = 0$	$\text{return } v\%3 = 0$

To prove the equivalence, we should check by induction the equality of both boolean results. Taking $g(x) = \lambda x.x = 0$, $g'(x) = \lambda x.(x \bmod 3) = 0$, the attempt to use Lemma 2 fails:

$$[x \bmod 3]_{0,+ \bmod 3} = 0 \iff [x]_{0,+ \bmod 3 \bmod 3} = 0 \\ \forall x, A, A'. A = 0 \iff A' \bmod 3 = 0 \implies (A + x \bmod 3) \bmod 3 = 0 \iff (A' + x) \bmod 3 = 0$$

The counter-example is: $A = 1, A' = 2, x = 1$. The hypothesis $A = 0 \iff A' \bmod 3 = 0$ is satisfied, but $(A + x \bmod 3) \bmod 3 \neq 0 \iff (A' + x) \bmod 3 = 0$ ($((1 + (1\%3))\%3 = 2, (2 + 1)\%3 = 0$)).

Despite the fact that Lemma 2 did not cover Example 3, it still belongs to a sub-class of Agg^1 for which an equivalence test method exists.

► **Definition 3 (The Agg_{sync}^1 class).** Let there be two Agg^1 programs P_1, P_2 with equal signature, whose program terms are $g_i([\varphi_i]_{init_i, f_i})$ for $i = 1, 2$. We say that $\langle P_1, P_2 \rangle \in Agg_{sync}^1$ if:

$$FV(\varphi_1) = FV(\varphi_2) \tag{4}$$

$$\forall \bar{v}_1, \bar{v}_2. \exists \bar{v}'. f_0(f_0(init_0, \varphi_0(\bar{v}_1)), \varphi_0(\bar{v}_2)) = f_0(init_0, \varphi_0(\bar{v}')) \\ \wedge f_1(f_1(init_1, \varphi_1(\bar{v}_1)), \varphi_1(\bar{v}_2)) = f_1(init_1, \varphi_1(\bar{v}')) \tag{5}$$

The Agg_{sync}^1 class contains pairs of programs in which multiple applications of the *fold* function starting from the same initial value and on the same sequence of valuations can be reduced to a single application of it, and it can be done using the same valuation for both programs.

► **Theorem 3 (Agg_{sync}^1 is decidable).** Let P_1, P_2 such that $\langle P_1, P_2 \rangle \in Agg_{sync}^1$, with input RDDs \bar{r} . We denote $\Phi(P_i) = g_i([\varphi_i]_{init_i, f_i})$. Then, $\Phi(P_1) = \Phi(P_2)$ if and only if:

$$g_1(init_1) = g_2(init_2) \tag{6}$$

$$\forall \bar{v}, \bar{y}, A_{\varphi_1}, A_{\varphi_2}. (A_{\varphi_1} = f_1(init_1, \varphi_1(\bar{y})) \wedge A_{\varphi_2} = f_2(init_2, \varphi_2(\bar{y})) \wedge g_1(A_{\varphi_1}) = g_2(A_{\varphi_2})) \\ \implies g_1(f_1(A_{\varphi_1}, \varphi_1(\bar{v}))) = g_2(f_2(A_{\varphi_2}, \varphi_2(\bar{v}))) \tag{7}$$

Proof. Sound (if): We prove the equality $g_1([\varphi_1]_{init_1, f_1}) = g_2([\varphi_2]_{init_2, f_2})$ by induction on the size of the RDDs $\llbracket \varphi_1 \rrbracket, \llbracket \varphi_2 \rrbracket$, denoted n . For $n = 0$, $\llbracket \varphi_1 \rrbracket(\bar{r}) = \llbracket \varphi_2 \rrbracket(\bar{r}) = \emptyset$, thus $[\varphi_i]_{init_i, f_i} = init_i$ ($i = 1, 2$), and the equality follows from Equation (6). Assuming for n and proving for $n+1$: We let a sequence of intermediate values $A_{\varphi_i, k}$, ($i = 1, 2; k = 1, \dots, n+1$), for which we know in particular that $g_1(A_{\varphi_1, n}) = g_2(A_{\varphi_2, n})$, and we need to prove $g_1(A_{\varphi_1, n+1}) = g_2(A_{\varphi_2, n+1})$. We denote $A_{\varphi_i, 0} = init_i$, and then we have $A_{\varphi_i, k} = f_i(A_{\varphi_i, k-1}, \varphi_i(\bar{a}_k))$ ($k = 1, \dots, n+1$) for some \bar{a}_k . According to Equation (5), $A_{\varphi_i, 2} = f_i(A_{\varphi_i, 1}, \varphi_i(\bar{a}_2)) = f_i(f_i(init_i, \varphi_i(\bar{a}_1)), \varphi_i(\bar{a}_2))$ yields $\exists \bar{a}'_2. \bigwedge_{i=1,2} A_{\varphi_i, 2} = f_i(init_i, \varphi_i(\bar{a}'_2))$. We can thus use Equation (5) to prove by induction that $\exists \bar{a}'_k. \bigwedge_{i=1,2} A_{\varphi_i, k} = f_i(init_i, \varphi_i(\bar{a}'_k))$, and in particular $\exists \bar{a}'_n. \bigwedge_{i=1,2} A_{\varphi_i, n} = f_i(init_i, \varphi_i(\bar{a}'_n))$. By applying Equation (7) for $\bar{v} = \overline{a_{n+1}}, \bar{y} = \bar{a}'_n$, we get:

$$\begin{aligned} g_1(f_1(f_1(init_1, \varphi_1(\bar{y})), \varphi_1(\bar{v}))) &= g_2(f_2(f_2(init_2, \varphi_2(\bar{y})), \varphi_2(\bar{v}))) \implies \\ g_1(f_1(f_1(init_1, \varphi_1(\bar{a}'_n)), \varphi_1(\bar{v}))) &= g_2(f_2(f_2(init_2, \varphi_2(\bar{a}'_n)), \varphi_2(\bar{v}))) \implies \\ g_1(f_1(A_{\varphi_1, n}, \varphi_1(\overline{a_{n+1}}))) &= g_2(f_2(A_{\varphi_2, n}, \varphi_2(\overline{a_{n+1}}))) \implies \\ g_1(A_{\varphi_1, n+1}) &= g_2(A_{\varphi_2, n+1}) \end{aligned}$$

as required.

Complete (only if): Assume towards a contradiction that either Equations (6) and (7) are false. If the requirement of Equation (6) is not satisfied, yet the aggregates are equivalent, i.e.

$$g_1([\varphi_1]_{init_1, f_1}) = g_2([\varphi_2]_{init_2, f_2}) \wedge g_1(init_1) \neq g_2(init_2)$$

then we can get a contradiction by choosing all input RDDs to be empty. Thus, $[\varphi_1]_{init_1, f_1} = init_1 \wedge [\varphi_2]_{init_2, f_2} = init_2 \implies g_1(init_1) = g_2(init_2)$, contradiction. The conclusion is that Equation (6) is a necessary condition for equivalence. Therefore, we assume just Equation (7) is false. Let there be counter-examples \bar{v}, \bar{y} to Equation (6) (The A_{φ_i} are determined immediately), and let:

$$F_i = f_i(f_i(init_i, \varphi_i(\bar{y})), \varphi_i(\bar{v}))$$

Then $g_1(F_1) \neq g_2(F_2)$. By Equation (5) we can write F_i as: $F_i = f_i(init_i, \varphi_i(\bar{w}))$ for some \bar{w} . We take an RDD $R = \{\{\bar{w}; 1\}\}$. Then $\llbracket \varphi_j \rrbracket(R) = \{\{\varphi_j(\bar{w}); 1\}\}$, for which: $\llbracket [\varphi_j]_{init_j, f_j} \rrbracket(R) = F_i$. By the assumption, $\llbracket g_1([\varphi_1]_{init_1, f_1}) \rrbracket(R) = \llbracket g_2([\varphi_2]_{init_2, f_2}) \rrbracket(R)$, but then $g_1(F_1) = g_2(F_2)$. Contradiction. \blacktriangleleft

► **Example 4 (Completing Example 3).** We have:

$$\begin{aligned} f_0(f_0(0, x \% 3), y \% 3) &= x \% 3 + y \% 3 \% 3 = f_0(0, (x + y) \% 3) = (x + y) \% 3 \% 3 \\ f_1(f_1(0, x), y) &= x + y = f_1(0, x + y) \end{aligned}$$

So Equation (5) is true (for arbitrary x, y , $x + y$ can reduce the two fold applications), and the programs belong to Agg_{sync}^1 . We are left with proving:

$$\forall x, y. ((0 + y \% 3) \% 3 = 0 \iff (0 + y) \% 3 = 0) \implies ((y + x \% 3) \% 3 = 0 \iff (y + x) \% 3 = 0)$$

which is correct — so we were able to prove the equivalence with the stronger lemma.

Note that checking if two programs P_1, P_2 belong to Agg_{sync}^1 involves a syntactic check of the free variables, and verification of an additional APA formula (Equation (5)).

► **Definition 4 (The Agg_R^1 class).** Let there be a program P with $\Phi(P) = \psi$. We say that $P \in Agg_R^1$ if ψ contains a single instance of an aggregate term $a = [\varphi]_{init, f}$. We write $\Phi(P) = \psi(a)$.

► **Lemma 4** (Lifting Lemma 2 to Agg_R^1). *Let there be two SparkLite programs $P_1, P_2 \in Agg_R^1$ with terms ψ_i and aggregate expressions $a_i = [\varphi_i]_{init_i, f_i}$, $i \in \{1, 2\}$. P_1 is equivalent to P_2 if:*

$$FV(\varphi_1) = FV(\varphi_2) \quad (8)$$

$$FV(\psi_1) = FV(\psi_2) \quad (9)$$

$$\forall \bar{x}. \psi_1(init_1)(\bar{x}) = \psi_2(init_2)(\bar{x}) \quad (10)$$

$$\forall \bar{v}, A_1, A_2. (\forall \bar{x}. \psi_1(A_1)(\bar{x}) = \psi_2(A_2)(\bar{x})) \implies (\forall \bar{x}. \psi_1(f_1(A_1, \varphi_1(\bar{v}))) (\bar{x}) = \psi_2(f_2(A_2, \varphi_2(\bar{v}))) (\bar{x})) \quad (11)$$

Lemmas 2,4 show that Agg^1, Agg_R^1 have a sound equivalence verification method, and Theorem 3 shows that Agg_{sync}^1 has a sound and complete equivalence verification method.⁸

5.3 Proof of undecidability of PE

We show a reduction of Hilbert's 10th problem to PE . We assume towards a contradiction that PE is decidable. Let there be a polynomial p over k variables x_1, \dots, x_k , and coefficients a_1, \dots, a_k . We use SparkLite operations and the input RDDs R_i to represent the value of the polynomial p for some valuation of the x_i . We define a translation φ from polynomials to SparkLite expressions:

- $\varphi(x_i) = [\text{map}(\lambda x.1)(R_i)]_{0, \lambda A, x. A+x}$
- $\varphi(x_i x_j) = [\text{cartesian}(\text{map}(\lambda x.1)(R_i), \text{map}(\lambda x.1)(R_j))]_{0, \lambda A, (x,y). A+x}$
- $\varphi(x_i^2) = [\text{cartesian}(\text{map}(\lambda x.1)(R_i), \text{map}(\lambda x.1)(R_i))]_{0, \lambda A, (x,y). A+x}$. This rule as well as the rest of the powers follows directly from the previous rule. For a degree k monom, we apply the *cartesian* operation k times, and fold it with $\lambda A, (x)_{i=1}^k. A + x_i$.
- $\varphi(x_i^0) = 1$, trivially.
- $\varphi(am(x_{i_1}, \dots, x_{i_j})) = a\varphi(m(x_{i_1}, \dots, x_{i_j}))$ where $m(x_{i_1}, \dots, x_{i_j}) = x_{i_1}^{l_1} \dots x_{i_j}^{l_j}$, for which φ is defined by induction according to the previous rules.
- $\varphi(p) = \sum_{i=1}^k \varphi(a_i m_i(x_{i_1}, \dots, x_{i_k}))$, where $p = \sum_{i=1}^k a_i m_i(x_{i_1}, \dots, x_{i_k})$.

Given a polynomial $p(a_1 \dots, a_k; x_1, \dots, x_k)$, we generate the following instance of the PE problem:

	$P1(R_1, \dots, R_k: RDD_{\text{Int}}):$	$P2(R_1, \dots, R_k: RDD_{\text{Int}}):$
1	return $\varphi(p) \neq 0$	return tt

By choosing input RDDs such that the size of R_i is equal to the matching variable x_i , we can simulate any valuation to the polynomial p . If $P1$ returns true, then the valuation is not a root of the polynomial p . Thus, if it is equivalent to the ‘true program’ $P2$, then the polynomial p has no roots. Therefore, if the algorithm solving PE outputs ‘equivalent’ then the polynomial p has no roots, and if it outputs ‘not equivalent’ then the polynomial p has some root, where $x_i = \llbracket R_i \rrbracket$ for the R_i ’s which are the witness for nonequivalence. Thus we have polynomial reduction to Hilbert's 10th problem.

SG: continue from here

⁸ Even when the completeness criterion for Agg_{sync}^1 is not met, we may be successful in proving the equivalence using Lemma 2. For example, $[((\lambda x.1)(\mathbf{x}_{r_0}), (\lambda x.1)(\mathbf{x}_{r_1}))]_{0, \lambda A, (x,y). A+x+y} = [((\lambda x.1)(\mathbf{x}_{r_1}), (\lambda x.1)(\mathbf{x}_{r_0}))]_{0, \lambda A, (x,y). A+x+y}$ can be proved by induction, but for $f = \lambda A, (x, y). A+x+y$, $f(f(A, (1, 1)), (1, 1)) = A+4 \neq f(A, (1, 1)) = A+2$ (the choice of valuation does not change the result). Thus, it does not satisfy the completeness criterion.

5.4 Multiple aggregates

5.4.0.1 Independent multiple aggregations.

Another relatively simple case of SparkLite programs is when the program contains multiple aggregate operations, but they are independent of each other. Namely, the result of one aggregate operations is not used in the other one.

► **Definition 5** (The Agg^n class). *Let there be a program P with $\Phi(P) = g([t_1]_{i_1, f_1}, \dots, [t_n]_{i_n, f_n})$. $P \in Agg^n$ if t_1, \dots, t_n do not contain aggregate terms.*

► **Lemma 5.** *Let P_1, P_2 be two programs in Agg^n , such that $\Phi(P_i) = g_i(\overline{[\varphi_i]_{init_i, f_i}})$. We have $g_1(\overline{[\varphi_1]_{init_1, f_1}}) = g_2(\overline{[\varphi_2]_{init_2, f_2}})$ if:*

$$\bigcup FV(\overline{\varphi_1}) = \bigcup FV(\overline{\varphi_2}) \quad (12)$$

$$g_1(\overline{init_1}) = g_2(\overline{init_2}) \quad (13)$$

$$\forall \bar{v}, \overline{A_{\varphi_1}}, \overline{A_{\varphi_2}}. g_1(\overline{A_{\varphi_1}}) = g_2(\overline{A_{\varphi_2}}) \implies g_1(\overline{f_1(A_{\varphi_1}, \varphi_1(\bar{v}))}) = g_2(\overline{f_2(A_{\varphi_2}, \varphi_2(\bar{v}))}) \quad (14)$$

► **Example 5** (Independent fold). Below are 2 programs which return a tuple containing the sum of positive elements in its first element, and the sum of negative elements in the second element. We show that by applying lemma 5, we are able to show the equivalence.

Let: $h : (\lambda(P, N), x.ite(x \geq 0, (P + x, N), (P, N - x)))$		
	P1($R: RDD_{Int}$):	P2($R: RDD_{Int}$):
1	return fold((0, 0), h)(R)	$R_P = \text{filter}(\lambda x. x \geq 0)(R)$
2		$R_N = \text{map}(\lambda x. -x)(\text{filter}(\lambda x. x < 0)(R))$
3		$p = \text{fold}(0, \lambda A, x. A + x)(R_P)$
4		$n = -\text{fold}(0, \lambda A, x. A + x)(R_N)$
5		return (p, n)

$$\Phi(P1) = [\mathbf{x}_R]_{(0,0),h}; \quad \Phi(P2) = ([\phi_{P2}(R_P)]_{0,+}, -[\phi_{P2}(R_N)]_{0,+})$$

$$\phi_{P2}(R_P) = ite(\mathbf{x}_R \geq 0, \mathbf{x}_R, \perp); \quad \phi_{P2}(R_N) = ite(\mathbf{x}_R < 0, -\mathbf{x}_R, \perp)$$

We let $g = \lambda(x, y).(x, y)$ and $g' = \lambda(x, y).(x, -y)$. Induction base case is trivial. Induction step:

$$\forall x, A, B, C. p_1(A) = B \wedge p_2(A) = C \implies h(A, x) = (B + ite(x \geq 0, x, 0), C + ite(x < 0, -x, 0))$$

Substituting for A with B, C , we get that we need to prove:

$$ite(x \geq 0, (B + x, C), (B, C - x)) \stackrel{?}{=} (B + ite(x \geq 0, x, 0), C + ite(x < 0, -x, 0))$$

which is a formula in the Augmented Presburger Arithmetic, whose validity is decidable. ■

5.4.0.2 Nested aggregations.

In the following subsection we present more complex SparkLite programs, in which the value of an aggregate operation is used in later aggregations (i.e. ‘nested’ aggregations). We see that the inductive method is sound in handling those cases, and that under certain conditions, it is complete too.

► **Example 6** (Conditional summation). The following example takes the *sum* of all elements which are greater than the *count* of elements in an RDD.

Let: $f : (\lambda A, (a, b). A + b)$
 $+: \lambda A, x. A + x$

	P1($R: RDD_{\text{Int}}$):	P2($R: RDD_{\text{Int}}$):
1	$R' = \text{map}(\lambda x.(x, 2))(R)$	$R' = \text{map}(\lambda x.(x, 1))(R)$
2	$sz = \text{fold}(0, f)(R')$	$sz = \text{fold}(0, f)(R')$
3	$B = \text{filter}(\lambda x.x > sz)(R)$	$B = \text{filter}(\lambda x.x > 2 * sz)(R)$
3	$\text{return fold}(0, +)(B)$	$\text{return fold}(0, +)(B)$

The equivalence condition is:

$$\left[\begin{array}{cc} \mathbf{x}_R & \mathbf{x}_R > \phi_{P1}(sz) \\ \perp & \text{otherwise} \end{array} \right]_{0,+} = \left[\begin{array}{cc} \mathbf{x}_R & \mathbf{x}_R > 2 * \phi_{P2}(sz) \\ \perp & \text{otherwise} \end{array} \right]_{0,+}$$

Replacing sz, sz' we get:

$$\left[\begin{array}{cc} \mathbf{x}_R & \mathbf{x}_R > [(\mathbf{x}_R^{(1)}, 2)]_{0,f} \\ \perp & \text{otherwise} \end{array} \right]_{0,+} = \left[\begin{array}{cc} \mathbf{x}_R & \mathbf{x}_R > 2 * [(\mathbf{x}_R^{(1)}, 1)]_{0,f} \\ \perp & \text{otherwise} \end{array} \right]_{0,+}$$

After formally applying lemma 2 we get that the above is equivalent if and only if $\phi_{P1}(sz) = 2 * \phi_{P2}(sz)$, that is:

$$[(\mathbf{x}_R^{(1)}, 2)]_{0,f} = 2 * [(\mathbf{x}_R^{(1)}, 1)]_{0,f}$$

In this case, we set $g = \lambda x.x, g' = \lambda x.2 * x$, and get:

$$0 = g(0) = 2 * 0 \tag{1}$$

$$\begin{aligned} \forall x, A, A'. A = 2 * A' &\implies f(A, (x, 2)) = A + 2 \\ &= 2 * A' + 2 \\ &= 2 * (A' + 1) \\ &= f(A', (x, 1)) \end{aligned} \tag{2}$$

Proving the equivalence. ■

This property is reflected in the following lemma:

► **Lemma 6.** *Let there be two SparkLite programs P_1, P_2 containing two aggregated expressions $a_i = [\varphi_i]_{\text{init}_i, f_i}$ for $i \in \{1, 2\}$. Let P_1, P_2 contain return aggregated expressions, the terms of which depend on the previous aggregations $a_i: h_j([\psi_j(a_j)]_{\text{init}'_j, g_j})$. If:*

$$FV(\psi_1) \simeq FV(\psi_2), \text{ denoted } FV_\psi \tag{1}$$

$$h_1(\text{init}_1) = h_2(\text{init}_2) \tag{2}$$

$$\forall \bar{v}, A_1, A_2. h_1(A_1) = h_2(A_2) \implies \tag{3}$$

$$h_1(g_1(A_1, \psi_1(a_1)[\bar{v}/FV_\psi])) = h_2(g_2(A_2, \psi_2(a_2)[\bar{v}/FV_\psi]))$$

Note that Lemma 6 is subject to all previously defined completeness criteria, in conjunction with the completeness criteria applied to the equivalence formula in the induction step. This allows us to define an algorithm for verifying complex equivalences with aggregated queries. The idea is to apply nested inductive proofs (on the nested expressions) during the proof of the induction step of the outer aggregations.

Several examples are given in appendix ??.

References

- 1 Java. <http://java.net>. Accessed: 2016-07-19.
- 2 Python. <https://www.python.org/>. Accessed: 2016-07-19.
- 3 Scala. <http://www.scala-lang.org>. Accessed: 2016-07-19.
- 4 Serge Abiteboul, Richard Hull, and Victor Vianu. *Foundations of Databases*. Addison-Wesley, 1995. URL: <http://webdam.inria.fr/Alice/>.
- 5 Michael Armbrust, Reynold S. Xin, Cheng Lian, Yin Huai, Davies Liu, Joseph K. Bradley, Xiangrui Meng, Tomer Kaftan, Michael J. Franklin, Ali Ghodsi, and Matei Zaharia. Spark sql: Relational data processing in spark. In *Proceedings of the 2015 ACM SIGMOD International Conference on Management of Data*, SIGMOD '15, pages 1383–1394, New York, NY, USA, 2015. ACM. URL: <http://doi.acm.org/10.1145/2723372.2742797>, doi:10.1145/2723372.2742797.
- 6 Aaron R. Bradley and Zohar Manna. *The Calculus of Computation: Decision Procedures with Applications to Verification*. Springer-Verlag New York, Inc., Secaucus, NJ, USA, 2007.
- 7 Surajit Chaudhuri and Moshe Y. Vardi. Optimization of real conjunctive queries. In *Proceedings of the Twelfth ACM SIGACT-SIGMOD-SIGART Symposium on Principles of Database Systems*, PODS '93, pages 59–70, New York, NY, USA, 1993. ACM. URL: <http://doi.acm.org/10.1145/153850.153856>, doi:10.1145/153850.153856.
- 8 E. F. Codd. A relational model of data for large shared data banks. *Commun. ACM*, 13(6):377–387, 1970. URL: <http://doi.acm.org/10.1145/362384.362685>, doi:10.1145/362384.362685.
- 9 Sara Cohen. Containment of aggregate queries. *SIGMOD Rec.*, 34(1):77–85, March 2005. URL: <http://doi.acm.org/10.1145/1058150.1058170>, doi:10.1145/1058150.1058170.
- 10 Sara Cohen, Werner Nutt, and Alexander Serebrenik. Rewriting aggregate queries using views. In *Proceedings of the Eighteenth ACM SIGMOD-SIGACT-SIGART Symposium on Principles of Database Systems*, PODS '99, pages 155–166, New York, NY, USA, 1999. ACM. URL: <http://doi.acm.org/10.1145/303976.303992>, doi:10.1145/303976.303992.
- 11 David C Cooper. Theorem proving in arithmetic without multiplication. *Machine Intelligence*, 1972.
- 12 Ron Cytron, Jeanne Ferrante, Barry K. Rosen, Mark N. Wegman, and F. Kenneth Zadeck. Efficiently computing static single assignment form and the control dependence graph. *ACM Trans. Program. Lang. Syst.*, 13(4):451–490, October 1991. URL: <http://doi.acm.org/10.1145/115372.115320>, doi:10.1145/115372.115320.
- 13 Jeffrey Dean and Sanjay Ghemawat. Mapreduce: Simplified data processing on large clusters. In *Proceedings of the 6th Conference on Symposium on Operating Systems Design & Implementation - Volume 6*, OSDI'04, pages 10–10, Berkeley, CA, USA, 2004. USENIX Association. URL: <http://dl.acm.org/citation.cfm?id=1251254.1251264>.
- 14 Michael J. Fischer and Michael O. Rabin. Super-exponential complexity of presburger arithmetic. Technical report, Massachusetts Institute of Technology, Cambridge, MA, USA, 1974.
- 15 Todd J. Green. Containment of conjunctive queries on annotated relations. In *Proceedings of the 12th International Conference on Database Theory*, ICDT '09, pages 296–309, New York, NY, USA, 2009. ACM. URL: <http://doi.acm.org/10.1145/1514894.1514930>, doi:10.1145/1514894.1514930.
- 16 Lauri Hella, Leonid Libkin, Juha Nurmonen, and Limsoon Wong. Logics with aggregate operators. *J. ACM*, 48(4):880–907, July 2001. URL: <http://doi.acm.org/10.1145/502090.502100>, doi:10.1145/502090.502100.

- 17 Ross Ihaka and Robert Gentleman. R: A language for data analysis and graphics. *Journal of Computational and Graphical Statistics*, 5(3):pp. 299–314, 1996.
- 18 Viktor Kuncak, Huu Hai Nguyen, and Martin C. Rinard. Deciding boolean algebra with presburger arithmetic. *J. Autom. Reasoning*, 36(3):213–239, 2006.
- 19 Aless Lasaruk and Thomas Sturm. *Effective Quantifier Elimination for Presburger Arithmetic with Infinity*, pages 195–212. Springer Berlin Heidelberg, Berlin, Heidelberg, 2009. URL: http://dx.doi.org/10.1007/978-3-642-04103-7_18, doi:10.1007/978-3-642-04103-7_18.
- 20 Derek C. Oppen. A 222pn upper bound on the complexity of presburger arithmetic. *Journal of Computer and System Sciences*, 16(3):323 – 332, 1978. URL: <http://www.sciencedirect.com/science/article/pii/0022000078900211>, doi:[http://dx.doi.org/10.1016/0022-0000\(78\)90021-1](http://dx.doi.org/10.1016/0022-0000(78)90021-1).
- 21 M. Presburger. Über die vollständigkeit eines gewissen systems der arithmetik ganzer zahlen, in welchem die addition als einzige operation hervortritt. *Comptes Rendus du I congrès de Mathématiciens des Pays Slaves*, pages 92–101, 1929.
- 22 Yehoshua Sagiv and Mihalis Yannakakis. Equivalences among relational expressions with the union and difference operators. *J. ACM*, 27(4):633–655, October 1980. URL: <http://doi.acm.org/10.1145/322217.322221>, doi:10.1145/322217.322221.
- 23 Matei Zaharia, Mosharaf Chowdhury, Tathagata Das, Ankur Dave, Justin Ma, Murphy McCauly, Michael J. Franklin, Scott Shenker, and Ion Stoica. Resilient distributed datasets: A fault-tolerant abstraction for in-memory cluster computing. In *Presented as part of the 9th USENIX Symposium on Networked Systems Design and Implementation (NSDI 12)*, pages 15–28, San Jose, CA, 2012. USENIX.
- 24 Matei Zaharia, Mosharaf Chowdhury, Michael J. Franklin, Scott Shenker, and Ion Stoica. Spark: Cluster computing with working sets. In *Proceedings of the 2Nd USENIX Conference on Hot Topics in Cloud Computing*, HotCloud’10, pages 10–10, Berkeley, CA, USA, 2010. USENIX Association. URL: <http://dl.acm.org/citation.cfm?id=1863103.1863113>.

A Extending Cooper's Algorithm to the Augmented Presburger Arithmetic

► **Proposition 5.** *The theory of formulas over \mathbb{Z}^n with terms in the Augmented Presburger Arithmetic is decidable.*

Proof. Let φ be a quantified formula over $\bigcup_n \mathbb{Z}^n$ with terms in the Augmented Presburger Arithmetic. We shall translate φ to a formula in the Presburger Arithmetic. For any atom $A: a = b$, and $a, b \in \mathbb{Z}^k$ for some $k > 0$, we build the following formula: $\bigwedge_{i=1}^k p_i(a) = p_i(b)$ and replace it in place of A . In the resulting formula, we assign new variable names, replacing the projected tuple variables: For $a \in \mathbb{Z}^k$ we define $x_{a,i} = p_i(a)$ for $i \in \{1, \dots, k\}$. Variable quantification extends naturally, i.e. $\forall a$ becomes $\forall x_{a,1}, \dots, x_{a,k}$, and similarly for \exists . ◀

B Typing rules for SparkLite

Booleans	$\frac{}{\rho \vdash \text{true} : \text{Boolean}}$	$\frac{}{\rho \vdash \text{false} : \text{Boolean}}$
Integers	$\frac{}{\rho \vdash 0, 1, \dots : \text{Integer}}$	
Integer ops	$\frac{\rho \vdash i : \text{Integer}, j : \text{Integer}, \text{op} \in \{+, -, *, \%\}}{\rho \vdash i \text{ op } j : \text{Integer}}$	$\frac{\rho \vdash i : \text{Integer}, j : \text{Integer}, \text{op} \in \{<, \leq, =, \geq, >\}}{\rho \vdash i \text{ op } j : \text{Boolean}}$
Boolean ops	$\frac{\rho \vdash b : \text{Boolean}}{\rho \vdash !b : \text{Boolean}}$	$\frac{\rho \vdash b_1 : \text{Boolean}, b_2 : \text{Boolean}, \text{op} \in \{\wedge, \vee\}}{\rho \vdash b_1 \text{ op } b_2 : \text{Boolean}}$
Tuples	$\frac{\rho \vdash e_1 : \tau_1, e_2 : \tau_2}{\rho \vdash (e_1, e_2) : \tau_1 \times \tau_2}$	$\frac{\rho \vdash e : \tau_1 \times \dots \times \tau_n}{\rho \vdash p_i(e) : \tau_i}$
UDFs	$\frac{\rho \vdash f : C_1 \times \dots \times C_n \rightarrow (\tau \rightarrow \tau'), \bar{e} : C_1 \times \dots \times C_n}{\rho \vdash f(\bar{e}) : \tau \rightarrow \tau'}$	$\frac{\rho \vdash f : \tau \rightarrow \tau', t : \tau}{\rho \vdash f(t) : \tau'}$
RDD	$\frac{\rho \vdash r : \text{RDD}_{\tau}, f : \tau \rightarrow \tau'}{\rho \vdash \text{map}(f)(r) : \text{RDD}_{\tau'}}$	$\frac{\rho \vdash r : \text{RDD}_{\tau}, f : \tau \rightarrow \text{Boolean}}{\rho \vdash \text{filter}(f)(r) : \text{RDD}_{\tau}}$
	$\frac{\rho \vdash r : \text{RDD}_{\tau}, r' : \text{RDD}_{\tau'}}{\rho \vdash \text{cartesian}(r, r') : \text{RDD}_{\tau \times \tau'}}$	$\frac{\rho \vdash r : \text{RDD}_{\tau}, f : \tau' \times \tau \rightarrow \tau, \text{init} : \tau'}{\rho \vdash \text{fold}(\text{init}, f)(r) : \tau'}$

■ **Figure 8** Typing rules for SparkLite

C Proof of Proposition 3

w.l.o.g. we $t_1 \neq \perp$ (symmetry). Then $\exists \bar{x}, y. y = t_1(\bar{x}) \wedge y \neq \perp$. We choose input RDDs \bar{r} such that $r_i = \{\{x_i; n_i\}\}$. If $t_2(\bar{x}) \neq y$ then $\llbracket t_1 \rrbracket(\bar{r}) \neq \llbracket t_2 \rrbracket(\bar{r})$, as required. Otherwise, the multiplicity of y in $\llbracket t_1 \rrbracket$ is $\prod_{i, r_i \in FV(t_1)} n_i$, and in $\llbracket t_2 \rrbracket$ it is $\prod_{i, r_i \in FV(t_2)} n_i$. As $FV(t_1) \neq FV(t_2)$ and $n_i > 0$, the multiplicities are different, thus $\llbracket t_1 \rrbracket(\bar{r}) \neq \llbracket t_2 \rrbracket(\bar{r})$, as required. ■

D Proof of Lemma 2

Proof. First we recall the semantics of the `fold` operation on some RDD R , which is a bag. We choose an arbitrary element $a \in R$ and apply the fold function recursively on a and on R with a single instance of a removed. We then write a sequence of elements in the order they are chosen by `fold`: $\langle a_1, \dots, a_n \rangle$, where n is the sum of all multiplicities in

the bag R . We also know that a requirement of aggregating operations' UDFs is that they are *commutative*, so the order of elements chosen does not change the final result. We also extend f_i to $\xi_i \times (\sigma_i \cup \{\perp\})$ by setting $f_i(A, \perp) = A$ (\perp is defined to behave as the neutral element for f_i). To prove $g([\varphi_0]_{init_0, f_0}) = g'([\varphi_1]_{init_1, f_1})$, it is necessary to prove that

$$g(\llbracket \text{fold} \rrbracket(f_0, init_0)(R_0)) = g'(\llbracket \text{fold} \rrbracket(f_1, init_1)(R_1))$$

We set $A_{\varphi_j, 0} = init_j$ for $j \in \{0, 1\}$. Each element of R_0, R_1 is expressible by providing a concrete valuation to the free variables of φ_0, φ_1 , namely the vector \bar{v} . We choose an arbitrary sequence of valuations to \bar{v} , denoted $\langle \bar{a}_1, \dots, \bar{a}_n \rangle$, and plug them into the *fold* operation for both R_0, R_1 . The result is 2 sequences of *intermediate values* $\langle A_{\varphi_0, 1}, \dots, A_{\varphi_0, n} \rangle$ and $\langle A_{\varphi_1, 1}, \dots, A_{\varphi_1, n} \rangle$. We have that $A_{\varphi_j, i} = f_j(A_{\varphi_j, i-1}, \varphi_j[\bar{a}_i/FV])$ for $j \in \{0, 1\}$, from the semantics of *fold*. Our goal is to show $g(A_{\varphi_0, n}) = g'(A_{\varphi_1, n})$ for all n . We prove the equality by induction on the *size* of the sequence of possible valuations of \bar{v} , denoted n . In each step i , we show $g(A_{\varphi_0, i}) = g'(A_{\varphi_1, i})$.

Case $n = 0$: $R_0 = R_1 = \emptyset$, so $\llbracket \text{fold} \rrbracket(f_0, init_0)(R_0) = init_0$ and $\llbracket \text{fold} \rrbracket(f_1, init_1)(R_1) = init_1$. From Equation (2), $g(init_0) = g'(init_1)$, as required.

Case $n = i$, assuming correct for $n \leq i - 1$: By assumption, we know that the sequence of intermediate values up to $i - 1$ is equal up to application of g, g' , and specifically $g(A_{\varphi_0, i-1}) = g'(A_{\varphi_1, i-1})$. We are given the i 'th concrete valuation of \bar{v} , denoted \bar{a}_i . We need to show $A_{\varphi_0, i} = A_{\varphi_1, i}$, so we use the formula for calculating the next intermediate value:

$$\begin{aligned} A_{\varphi_0, i} &= f_0(A_{\varphi_0, i-1}, \varphi_0[\bar{a}_i/FV]) \\ A_{\varphi_1, i} &= f_1(A_{\varphi_1, i-1}, \varphi_1[\bar{a}_i/FV]) \end{aligned}$$

We use Equation (3), plugging in $\bar{v} = \bar{a}_i$, $A_{\varphi_0} = A_{\varphi_0, i-1}$, and $A_{\varphi_1} = A_{\varphi_1, i-1}$. By the induction assumption, $g(A_{\varphi_0, i-1}) = g'(A_{\varphi_1, i-1})$, therefore $g(A_{\varphi_0}) = g'(A_{\varphi_1})$, so Equation (3) yields $g(f_0(A_{\varphi_0}, \varphi_0[\bar{a}_i/FV])) = g'(f_1(A_{\varphi_1}, \varphi_1[\bar{a}_i/FV]))$. By substituting back A_{φ_j} and the formula for the next intermediate value, we get: $g(A_{\varphi_0, i}) = g'(A_{\varphi_1, i})$ as required. ◀